Sequence-to-Sequence Architectures

Kapil Thadani
kapil@cs.columbia.edu

YAHOO!
RESEARCH

Outline

- Machine translation
  - Phrase-based MT
  - Encoder-decoder architecture
- Attention mechanism
  - Bahdanau et al (2015)
  - Visualizations
  - Variants
  - Transformers
- Decoding large vocabularies
  - Alternative approaches
  - Copying mechanism
- Autoencoders
  - Denoising autoencoders
  - Variational autoencoders (VAEs)

Previously: processing text with RNNs

Inputs

- One-hot vectors for words/characters/previous output
- Embeddings for words/sentences/context
- CNN over characters/words/sentences

⋮

Recurrent layers

- Forward, backward, bidirectional, deep
- Activations: σ, tanh, gated (LSTM, GRU), ReLU initialized with identity

⋮

Outputs

- Softmax over words/characters/labels
- Absent (i.e., pure encoders)

⋮

Machine Translation

"*One naturally wonders if the problem of translation could conceivably be treated as a problem in cryptography. When I look at an article in Russian, I say: 'This is really written in English, but it has been coded in some strange symbols. I will now proceed to decode.*"

— Warren Weaver
Translation (1955)

The MT Pyramid

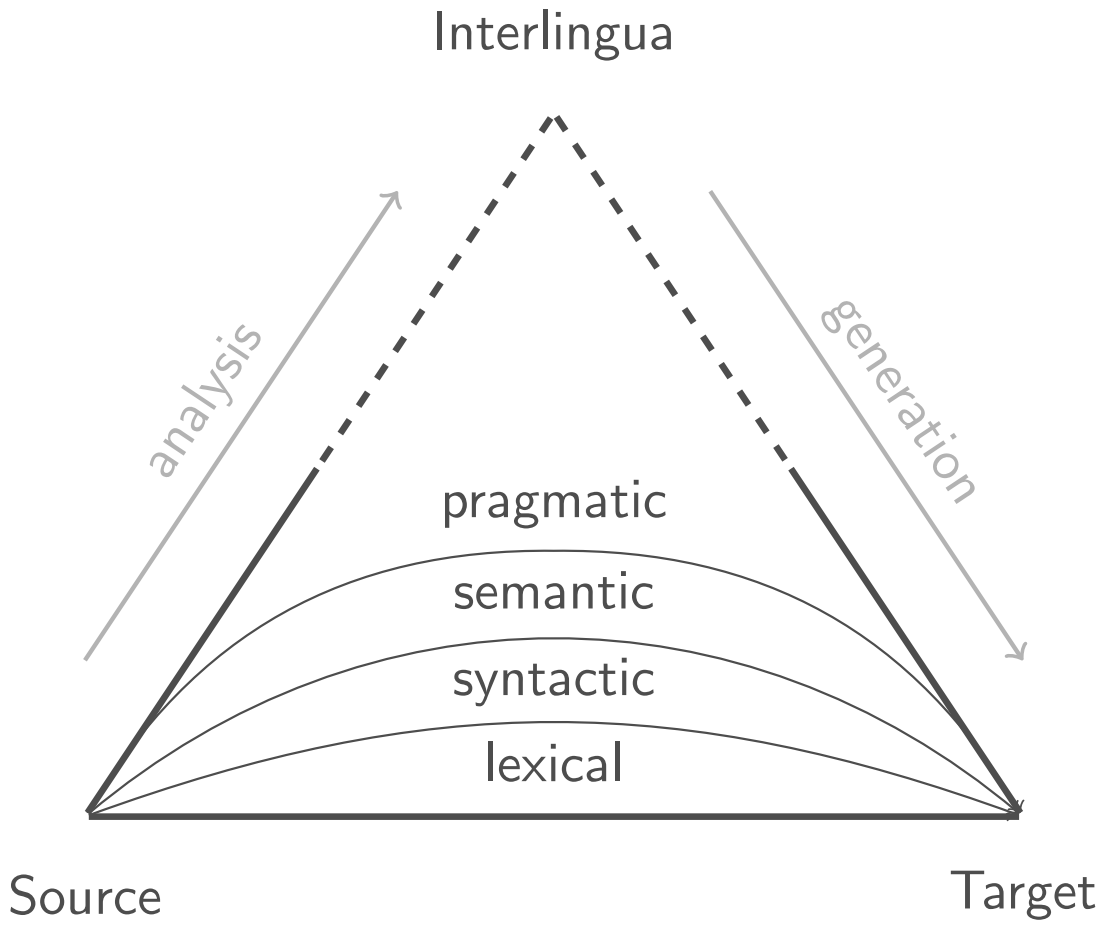

Phrase-based MT

Tomorrow I will fly to the conference in Canada

Morgen fliege Ich nach Kanada zur Konferenz

Phrase-based MT

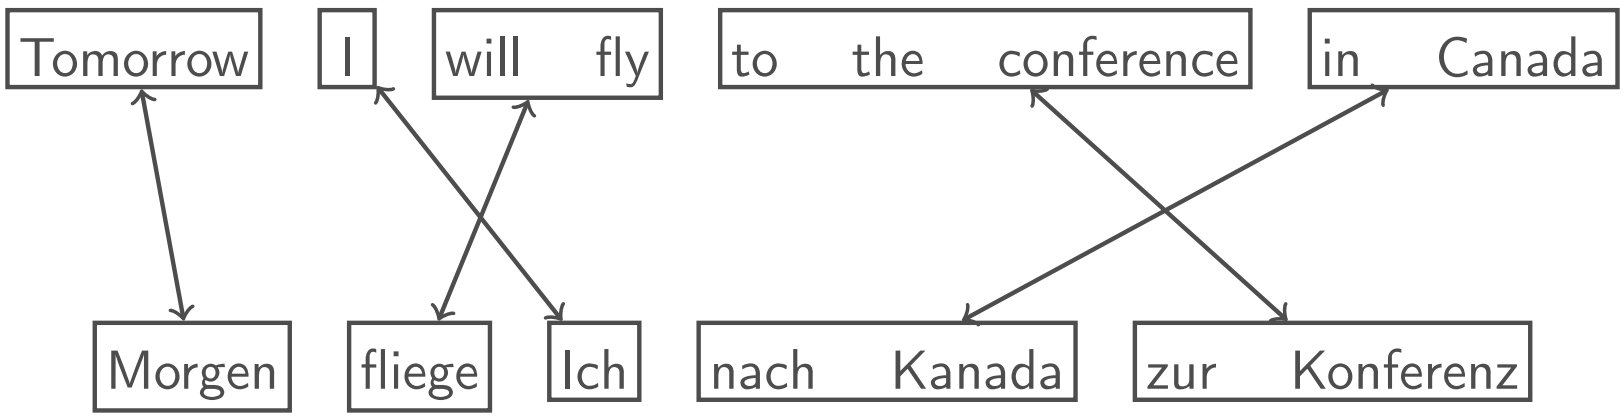

Phrase-based MT

1. Collect bilingual dataset $\langle S_i, T_i \rangle \in \mathcal{D}$

2. Unsupervised phrase-based alignment
   - phrase table π

3. Unsupervised n-gram language modeling
   - language model ψ

4. Supervised decoder
   - parameters θ

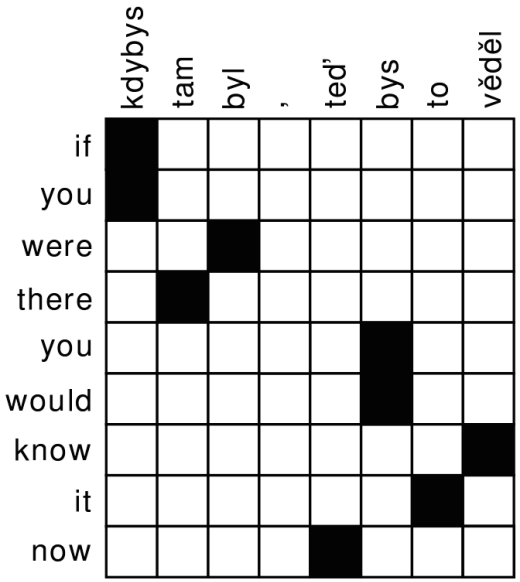

$$\widehat{T} = \arg\max_{T} \; p(T|S)$$

$$= \arg\max_{T} \; p(S|T, \pi, \theta) \cdot p(T|\psi)$$

Neural MT

1. Collect bilingual dataset $\langle S_i, T_i \rangle \in \mathcal{D}$
2. Unsupervised phrase-based alignment
   - phrase table π
3. Unsupervised n-gram language modeling
   - language model ψ
4. Supervised encoder-decoder framework
   - parameters θ

RNN

Input words $x_1, \ldots, x_n$

Output label z

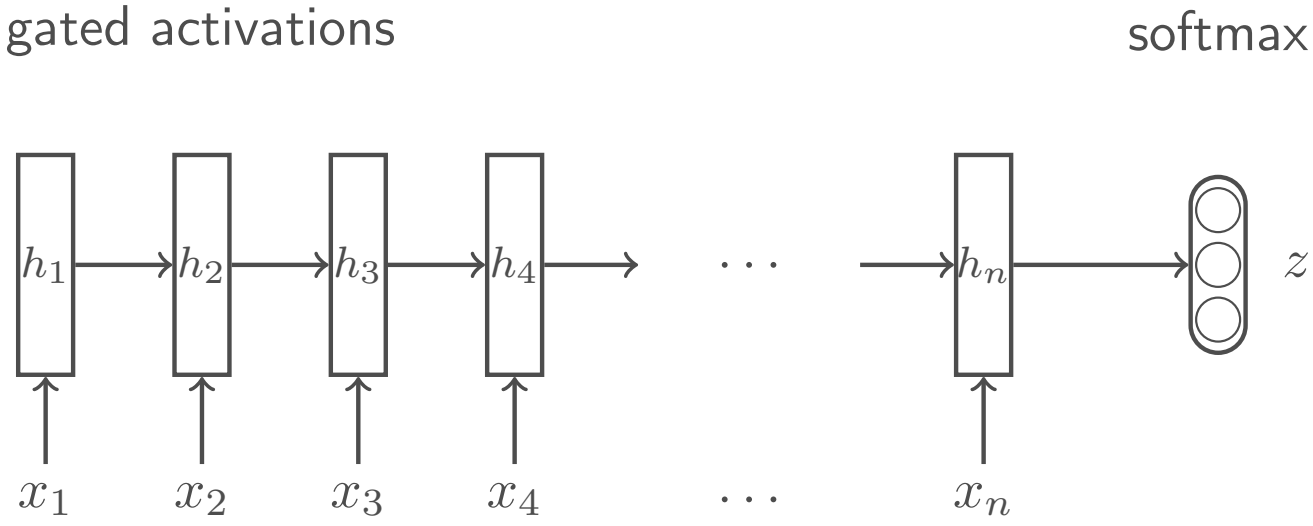

Deep RNN

Input words $x_1, \ldots, x_n$

Output label z

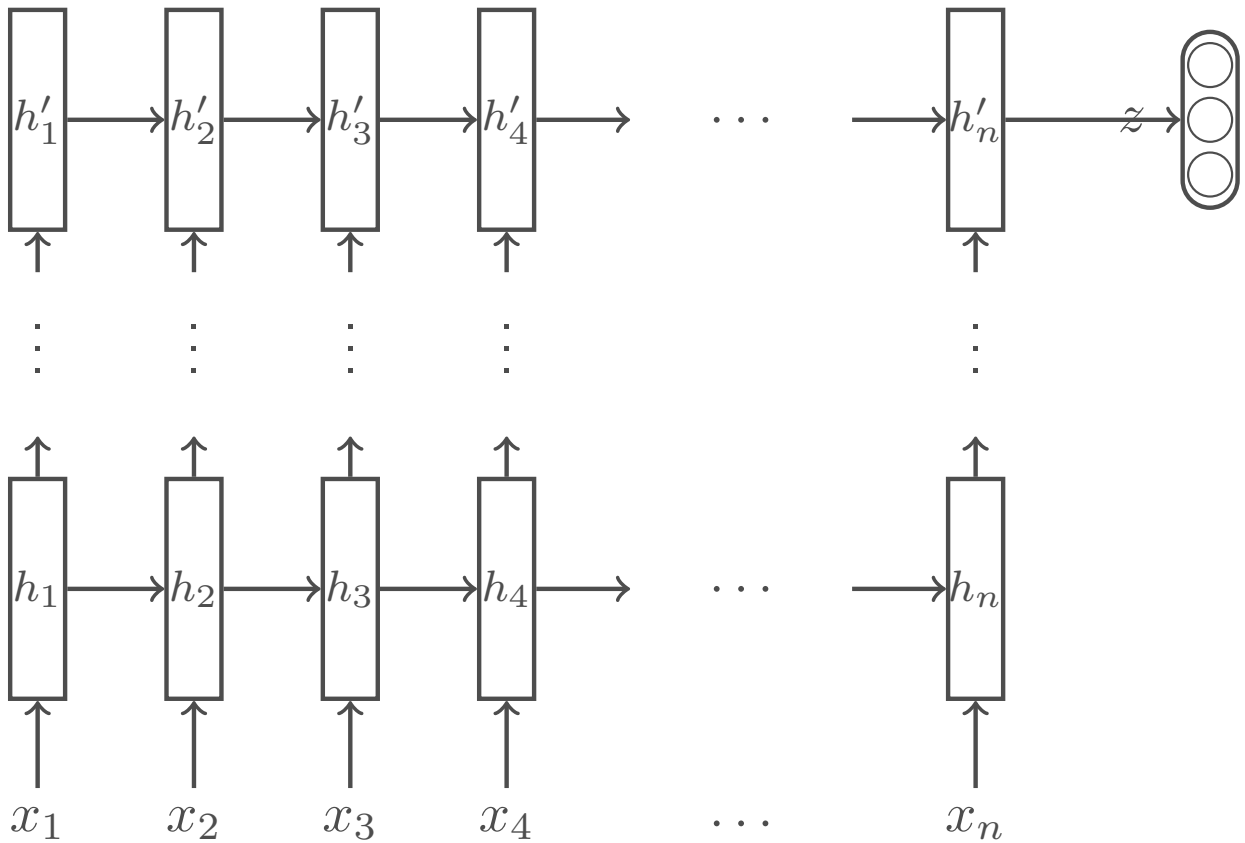

Bidirectional RNN

Input words $x_1, \ldots, x_n$

Output label z

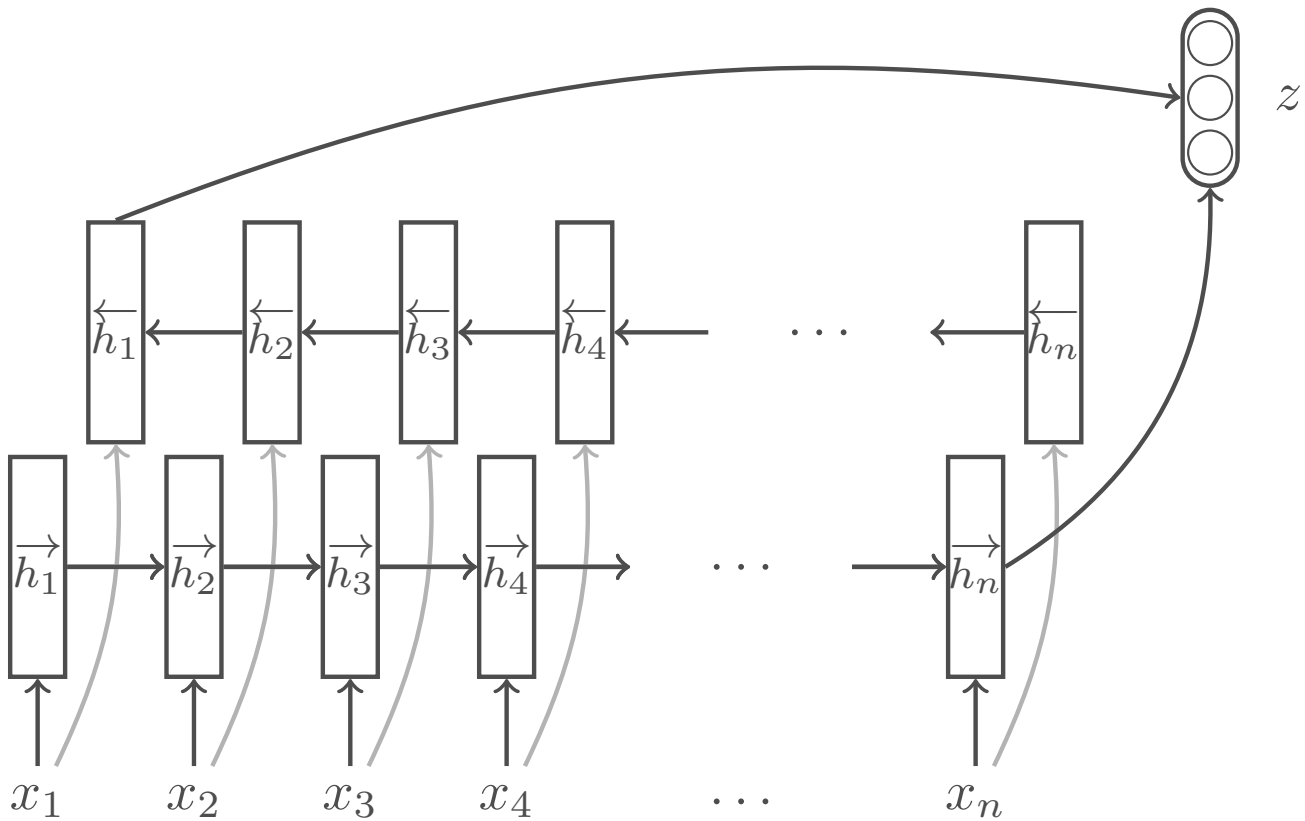

RNN encoder

Input words $x_1, \ldots, x_n$

Output encoding c

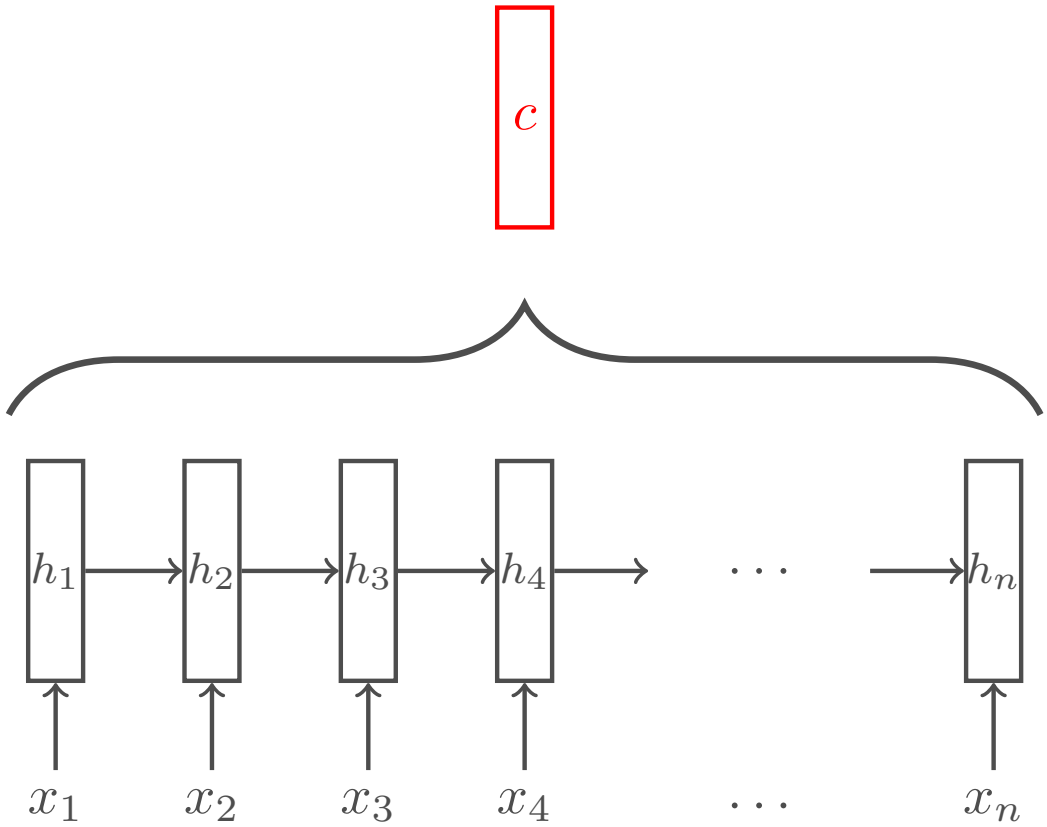

RNN language model

Input words $y_1, \ldots, y_k$

Output following words $y_k, \ldots, y_m$

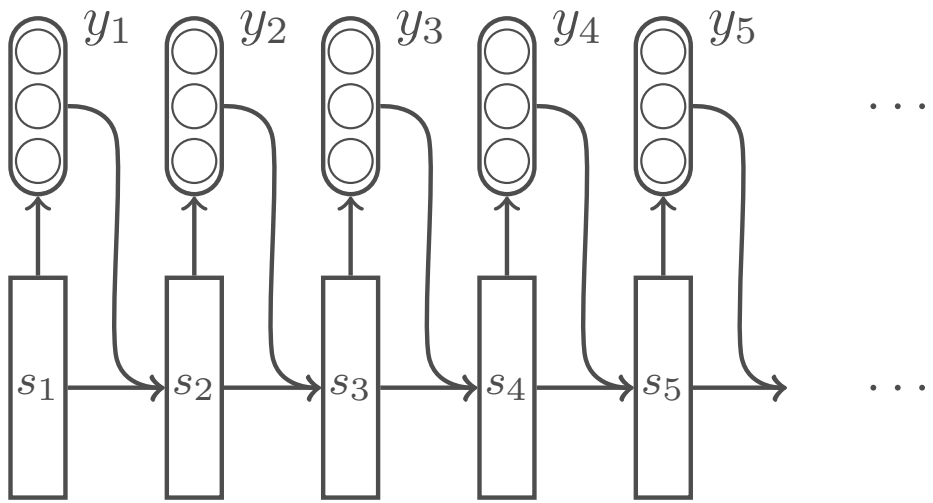

RNN decoder

Input context c

Output words $y_1, \ldots, y_m$

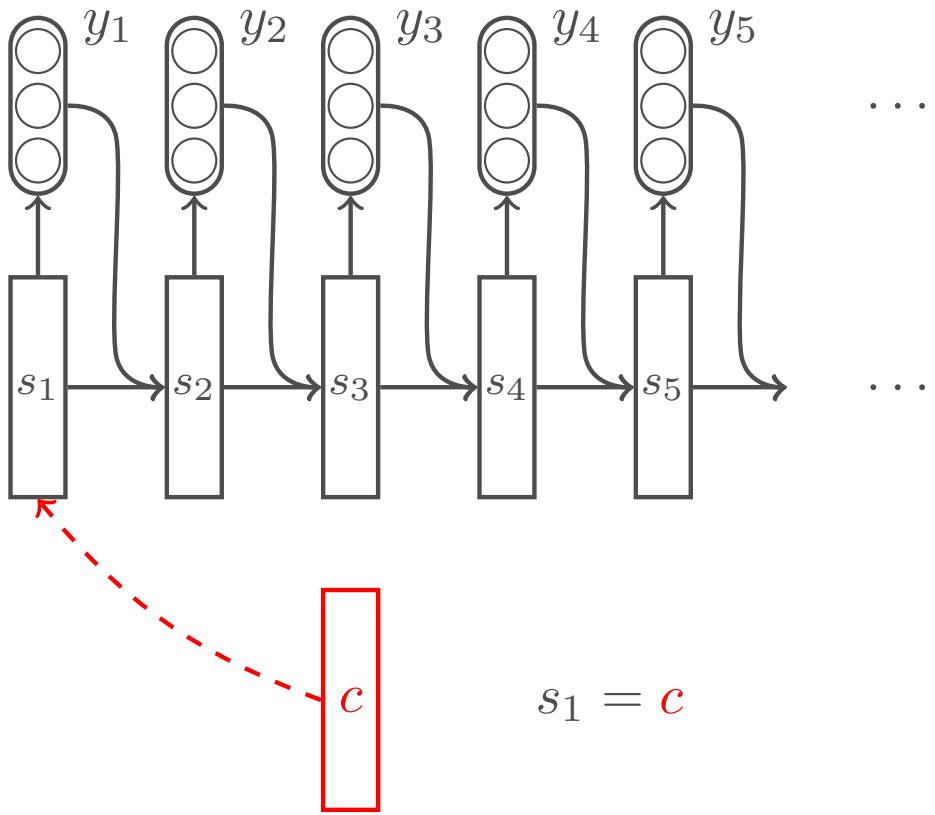

RNN decoder

Input context c

Output words $y_1, \ldots, y_m$

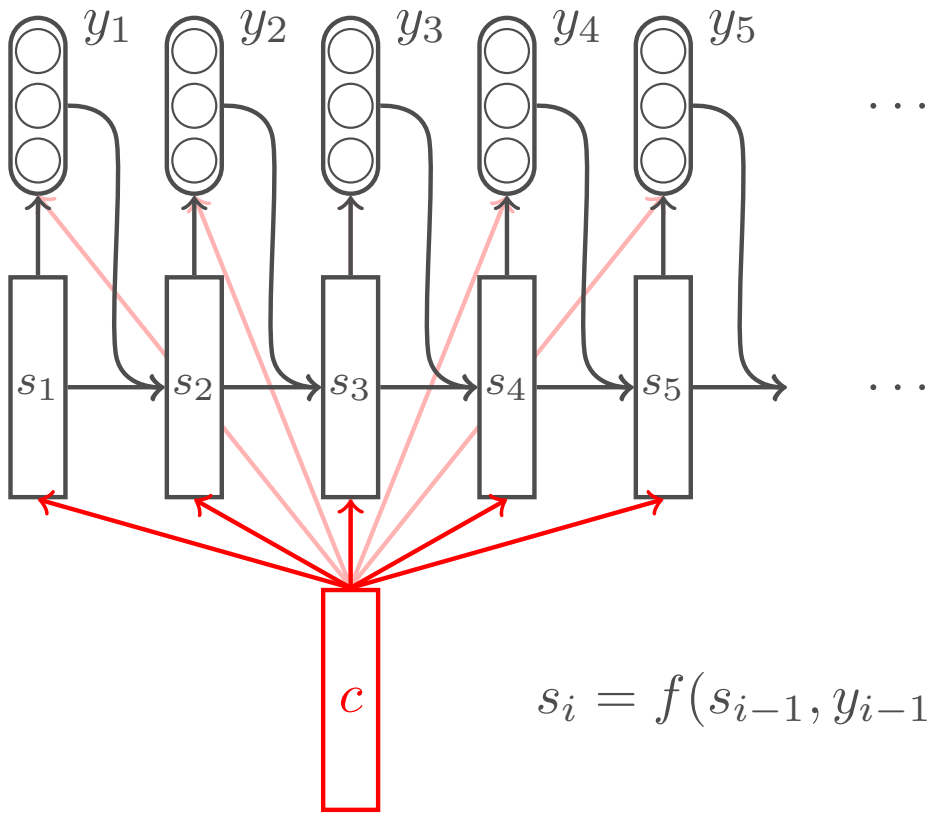

$$s_i = f(s_{i-1}, y_{i-1}, c)$$

Sequence-to-sequence learning

Sutskever, Vinyals & Le (2014)

Sequence to Sequence Learning with Neural Networks

Input words $x_1, \ldots, x_k$

Output words $y_1, \ldots, y_m$

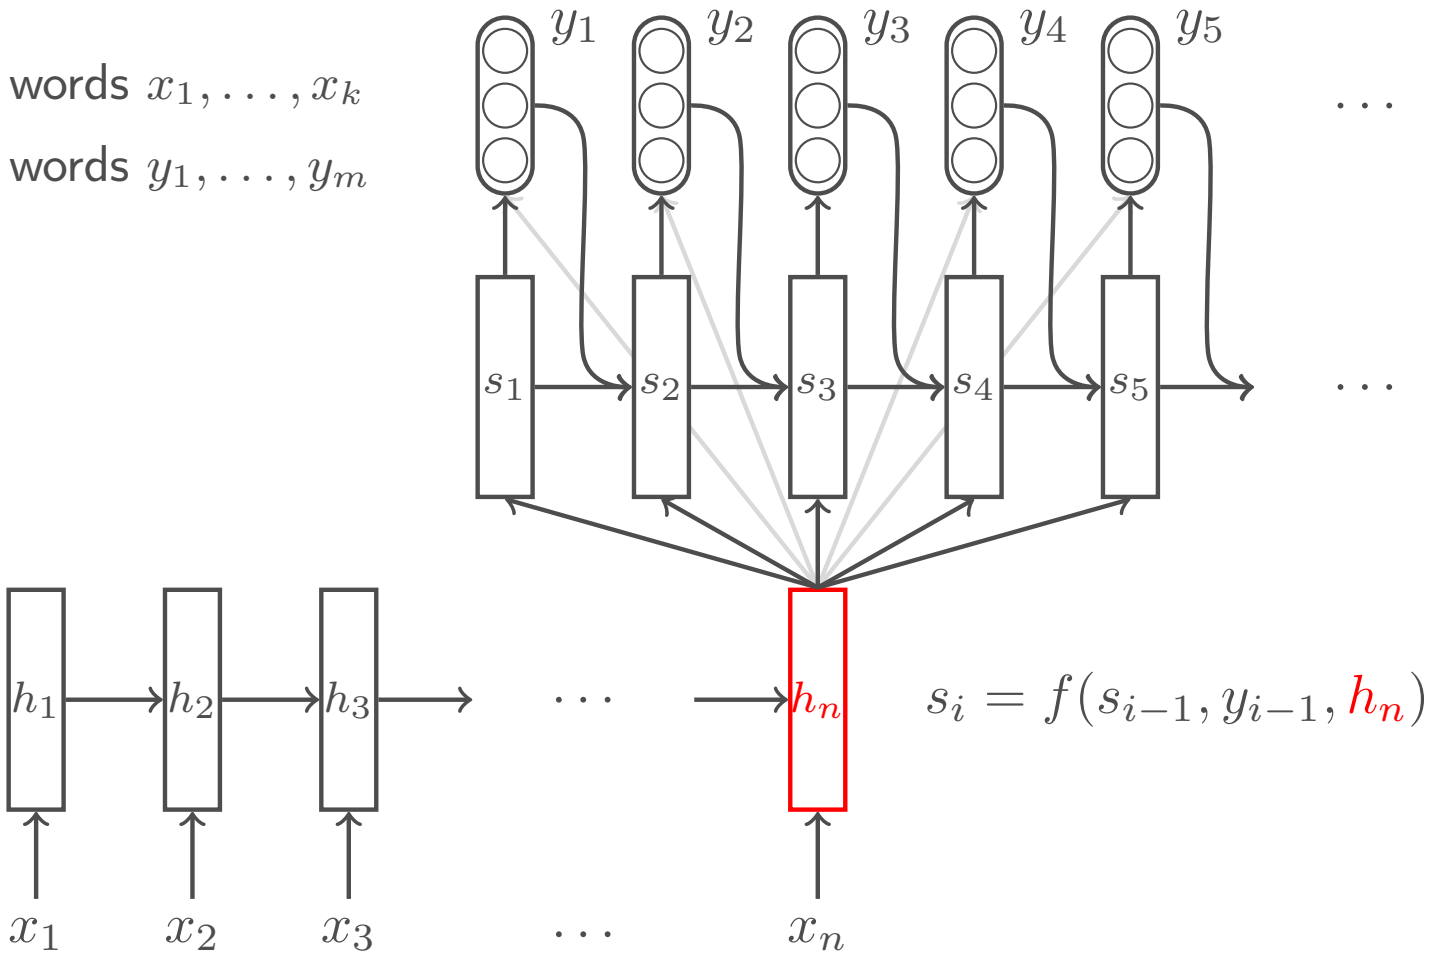

Sequence-to-sequence learning

Sutskever, Vinyals & Le (2014)

Sequence to Sequence Learning with Neural Networks

Produces a fixed length representation of input

- "sentence embedding" or "thought vector"

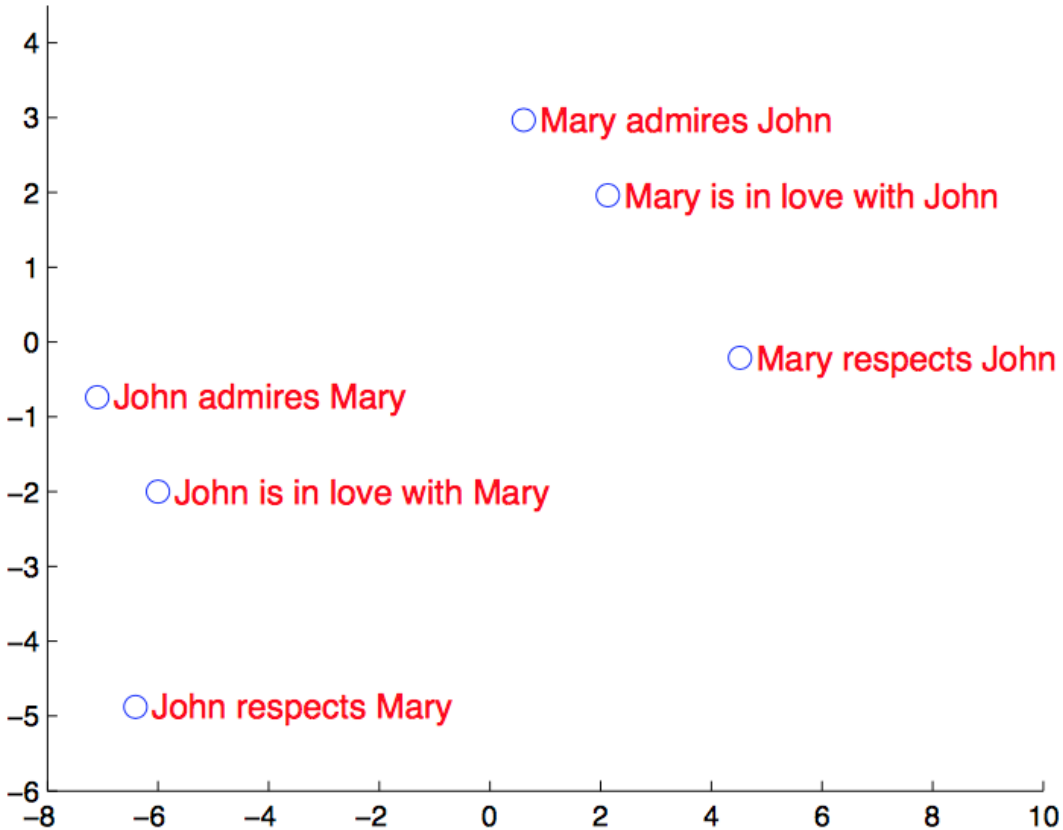

Sequence-to-sequence learning

Sutskever, Vinyals & Le (2014)

Sequence to Sequence Learning with Neural Networks

Produces a fixed length representation of input

- "sentence embedding" or "thought vector"

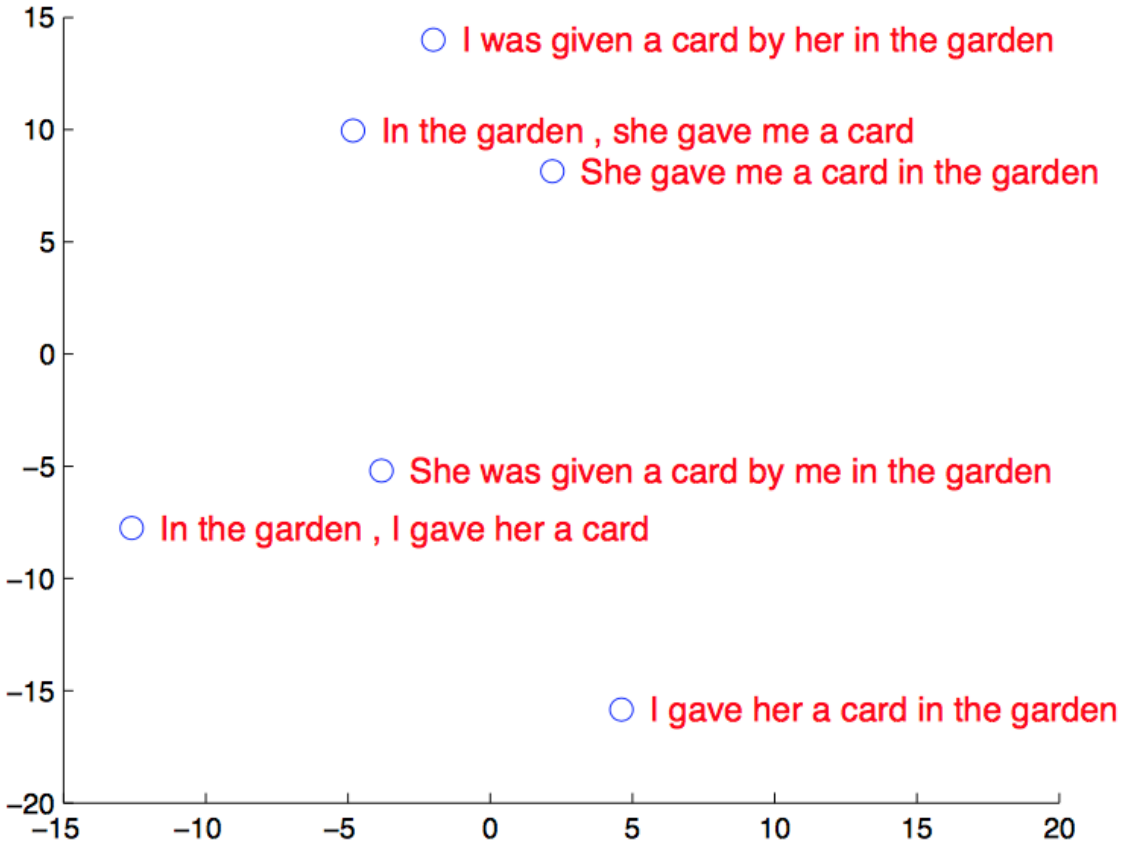

Sequence-to-sequence learning

Sutskever, Vinyals & Le (2014)

Sequence to Sequence Learning with Neural Networks

LSTM units do not solve vanishing gradients

- Poor performance on long sentences
- Need to reverse the input

Method	test BLEU score (ntst14)
Bahdanau et al. [2]	28.45
Baseline System [29]	33.30
Single forward LSTM, beam size 12	26.17
Single reversed LSTM, beam size 12	30.59

Attention-based translation

Bahdanau et al (2015)

Neural Machine Translation by Jointly Learning to Align and Translate

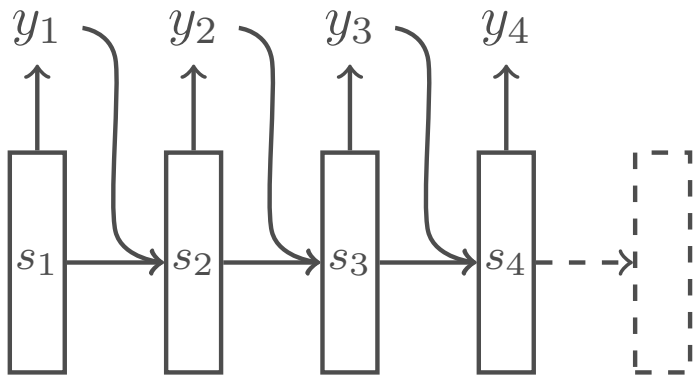

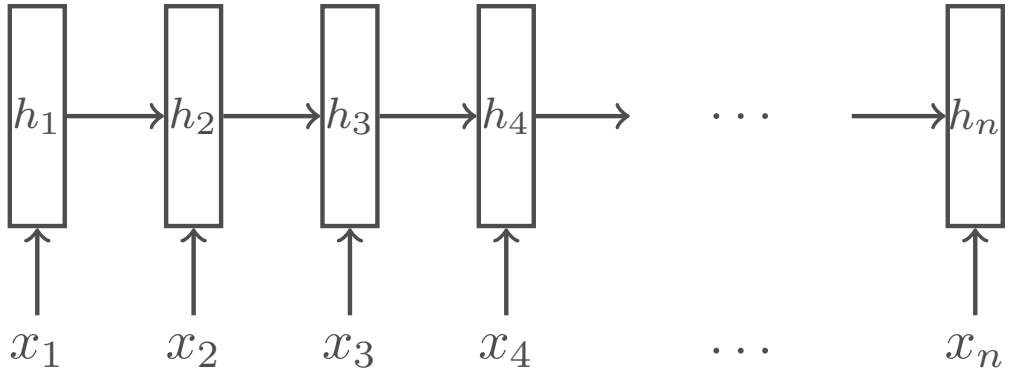

Attention-based translation

Bahdanau et al (2015)

Neural Machine Translation by Jointly Learning to Align and Translate

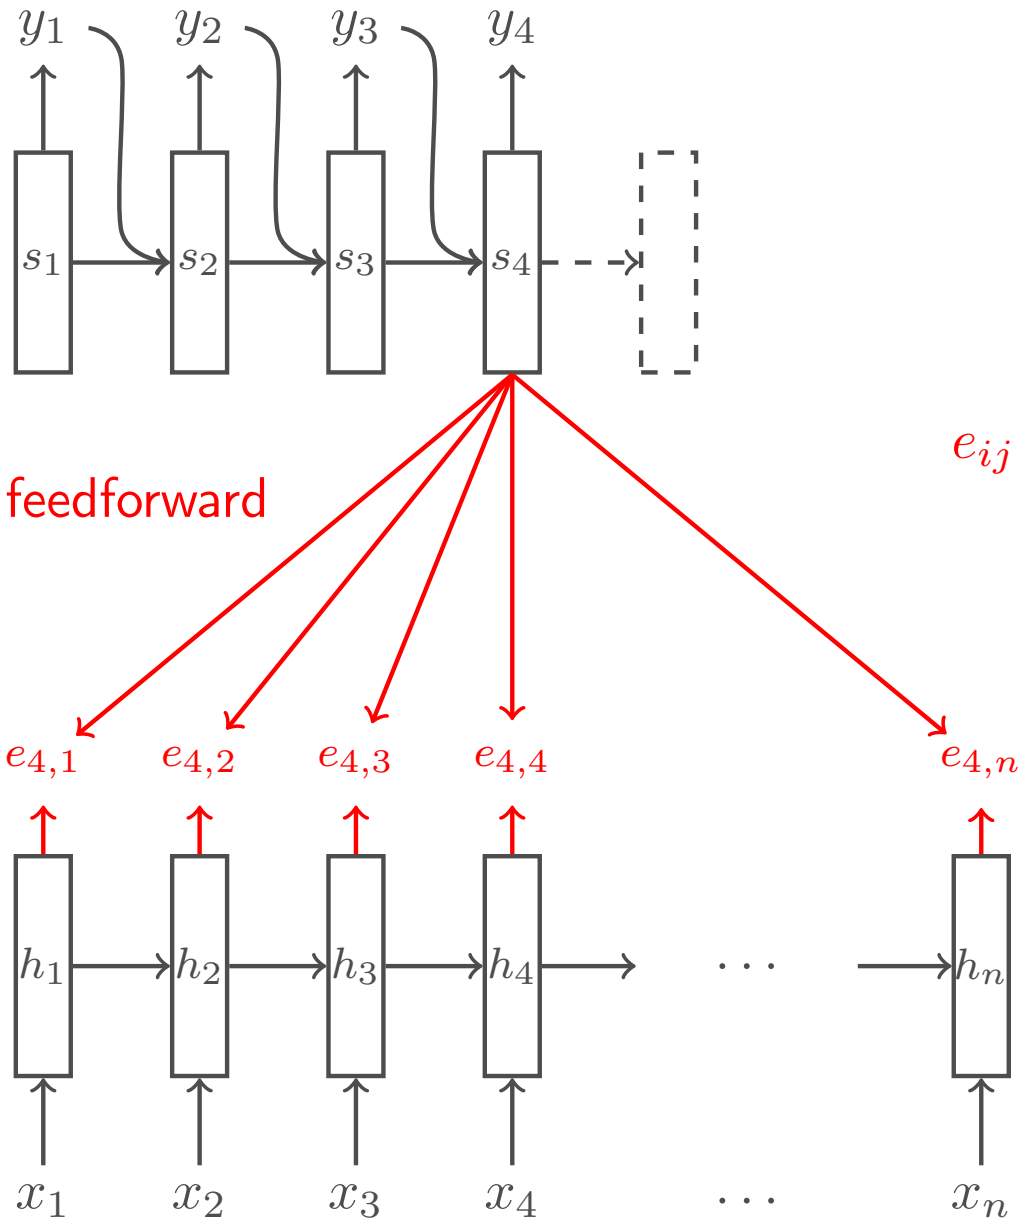

$$e_{ij} = a(s_{i-1}, h_j)$$

Attention-based translation

Bahdanau et al (2015)

Neural Machine Translation by Jointly Learning to Align and Translate

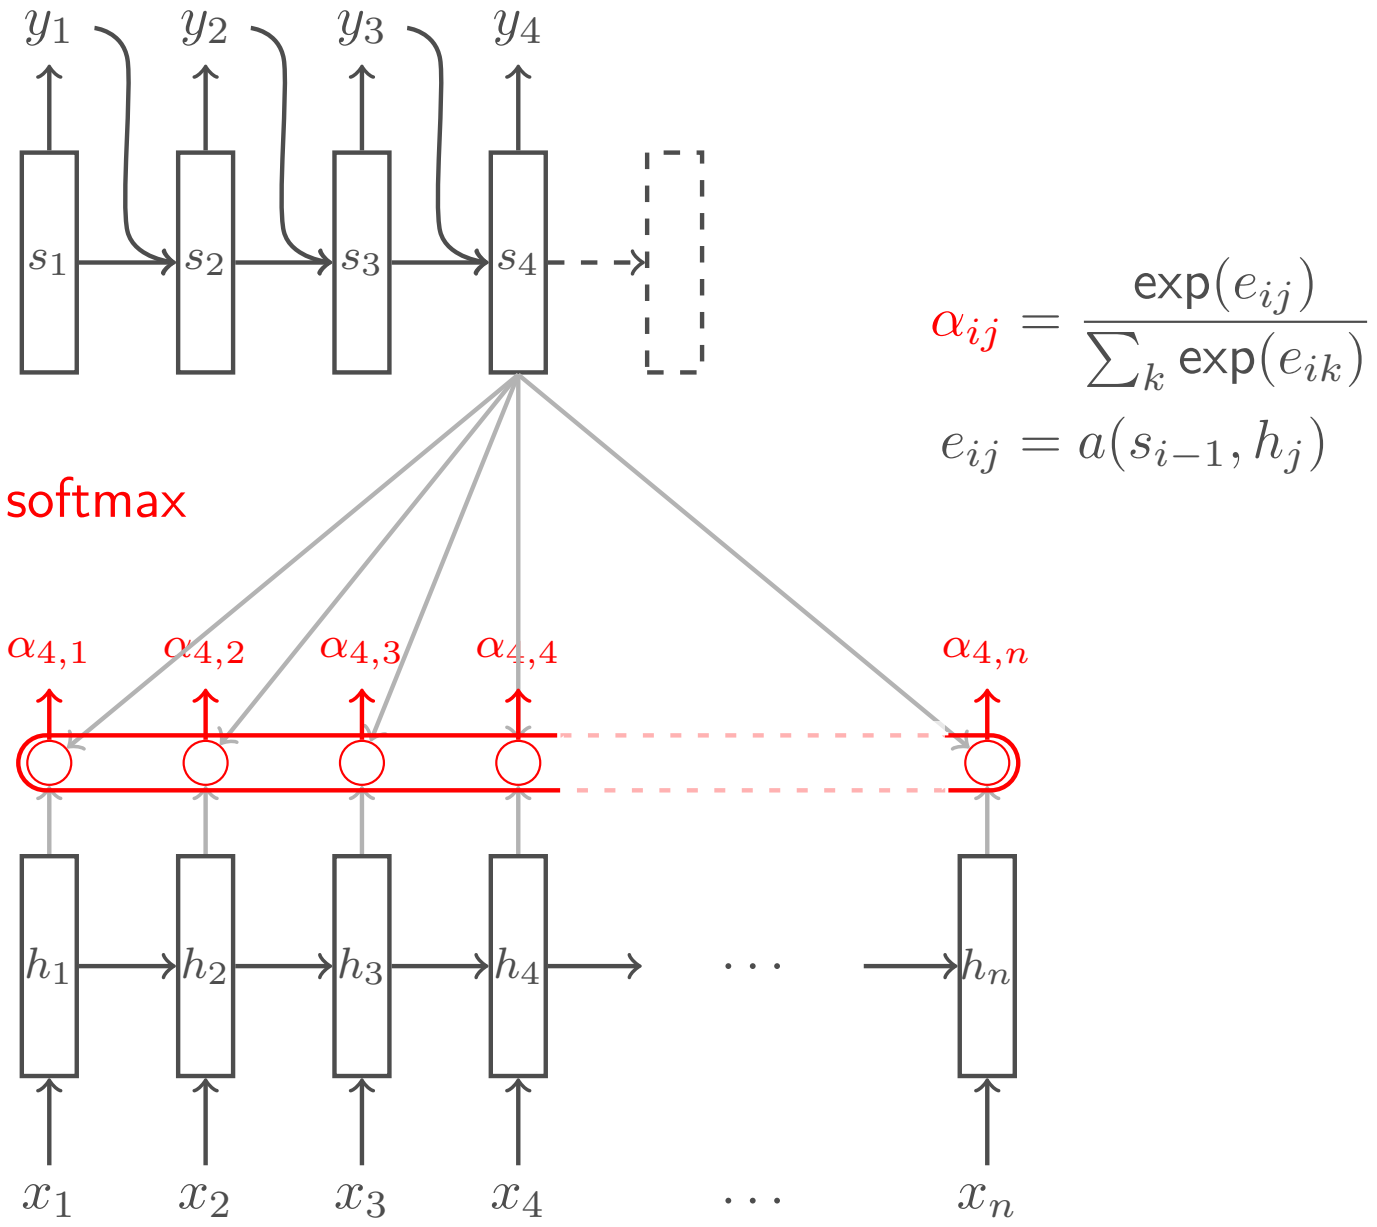

$$\alpha_{ij} = \frac{\exp(e_{ij})}{\sum_k \exp(e_{ik})}$$

$$e_{ij} = a(s_{i-1}, h_j)$$

Attention-based translation

Bahdanau et al (2015)

Neural Machine Translation by Jointly Learning to Align and Translate

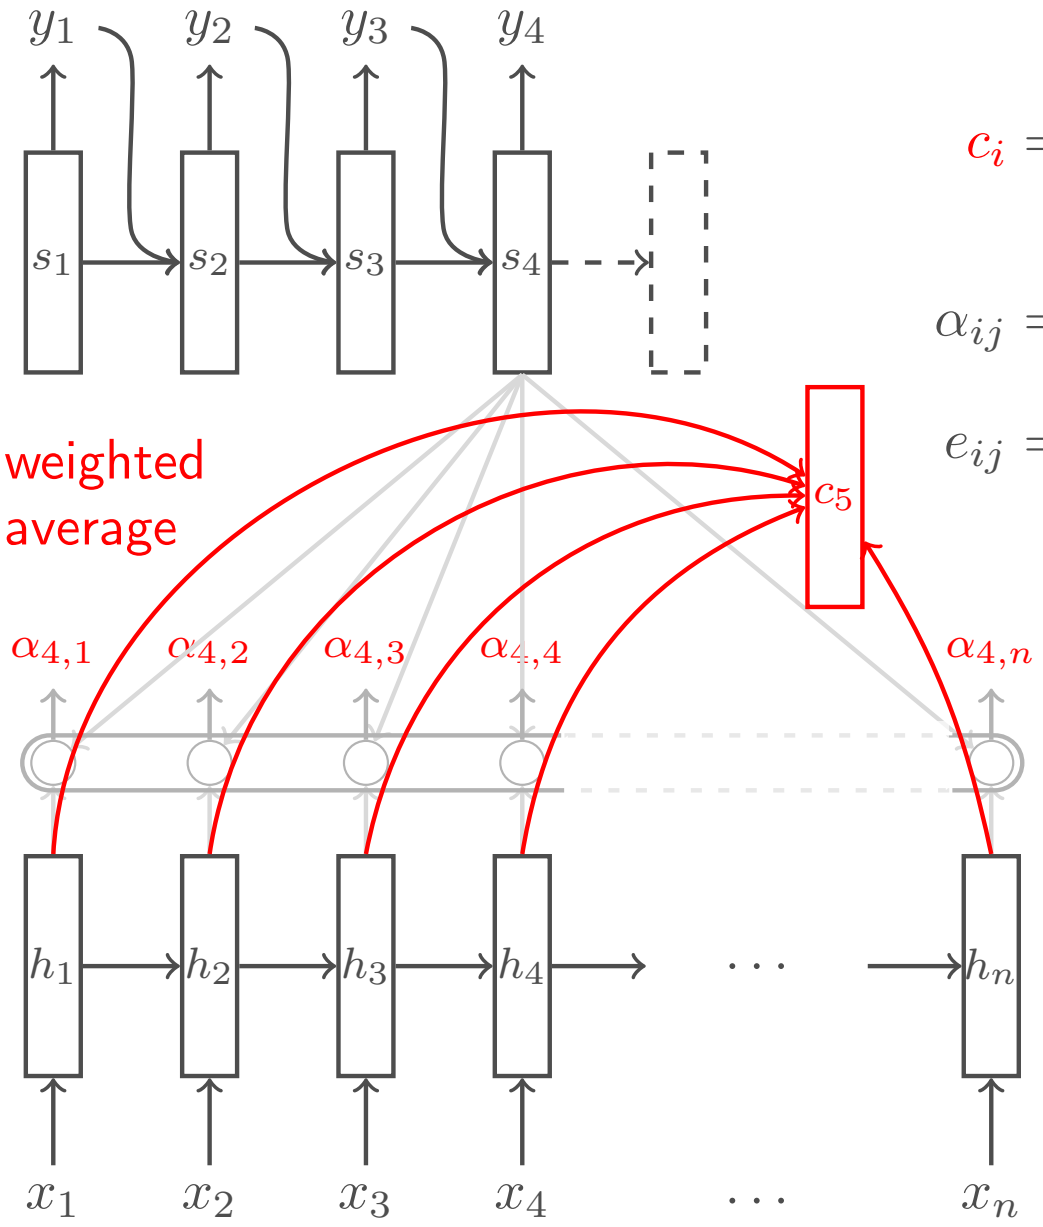

$$c_i = \sum_j \alpha_{ij} h_j$$

$$\alpha_{ij} = \frac{\exp(e_{ij})}{\sum_k \exp(e_{ik})}$$

$$e_{ij} = a(s_{i-1}, h_j)$$

Attention-based translation

Bahdanau et al (2015)

Neural Machine Translation by Jointly Learning to Align and Translate

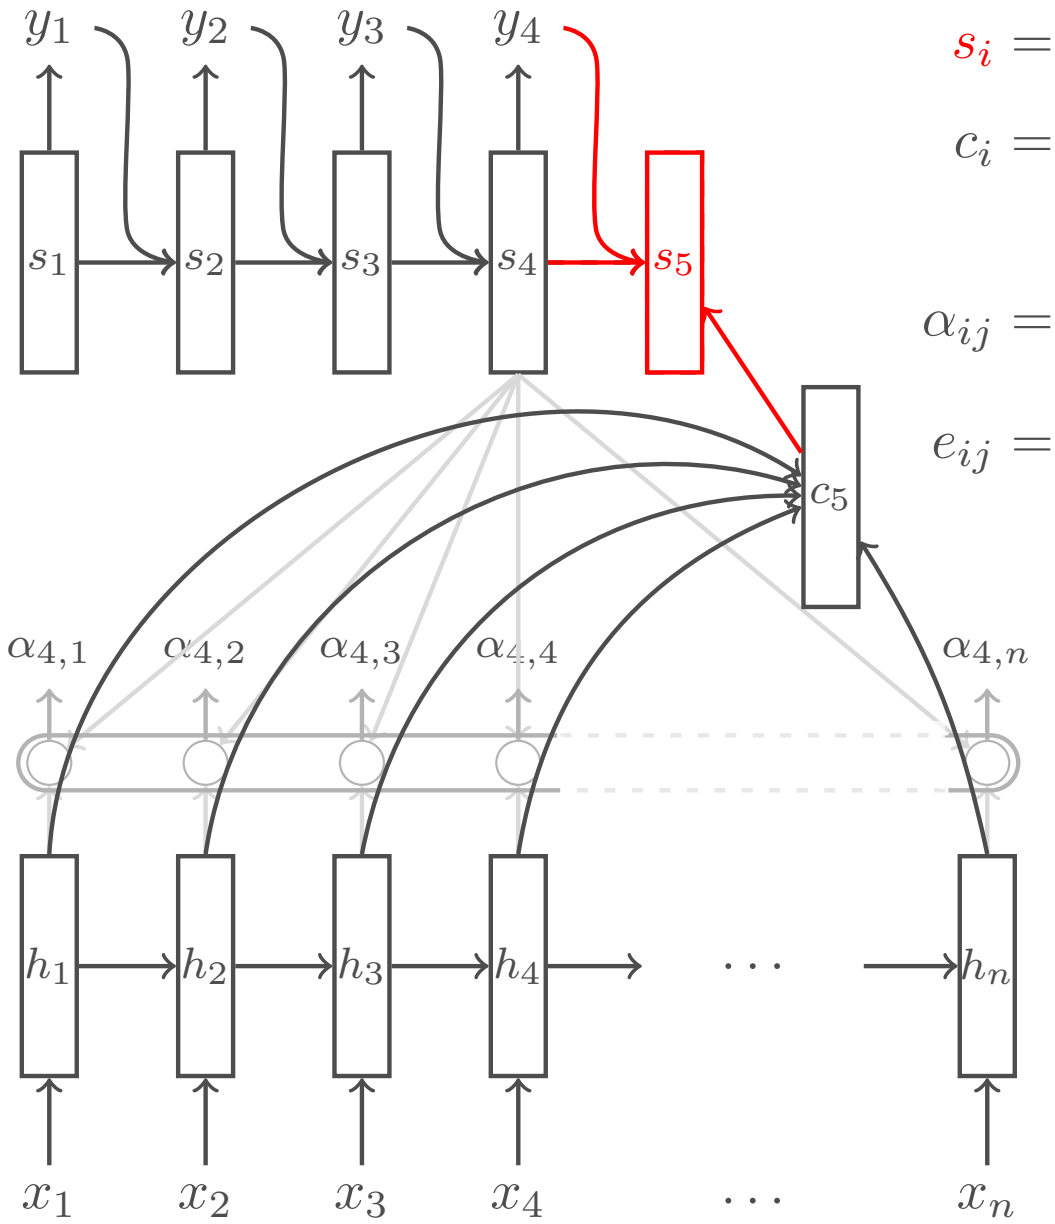

$$s_i = f(s_{i-1}, y_{i-1}, c_i)$$

$$c_i = \sum_j \alpha_{ij} h_j$$

$$\alpha_{ij} = \frac{\exp(e_{ij})}{\sum_k \exp(e_{ik})}$$

$$e_{ij} = a(s_{i-1}, h_j)$$

Attention-based translation

Bahdanau et al (2015)

Neural Machine Translation by Jointly Learning to Align and Translate

- Bidirectional encoder, GRU activations
- Softmax for y_i depends on y_{i-1} and an additional hidden layer

\+ Backprop directly to attended regions, avoiding vanishing gradients

\+ Can visualize attention weights α_{ij} to interpret prediction

– Inference is $\mathcal{O}(mn)$ instead of $\mathcal{O}(m)$ for seq-to-seq

Attention-based translation

Bahdanau et al (2015)

Neural Machine Translation by Jointly Learning to Align and Translate

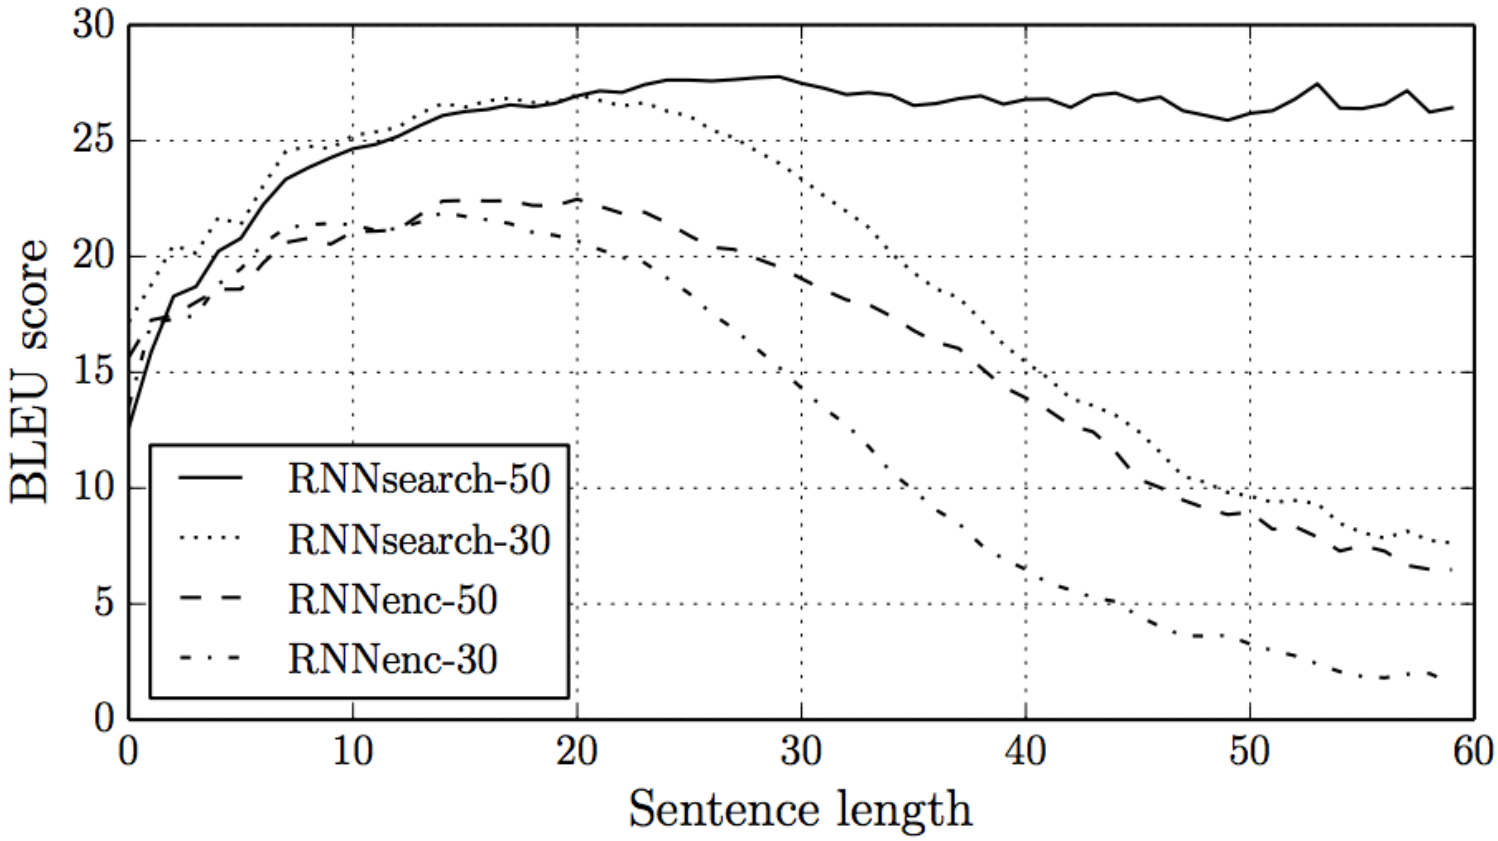

Improved results on long sentences

Attention-based translation

Bahdanau et al (2015)

Neural Machine Translation by Jointly Learning to Align and Translate

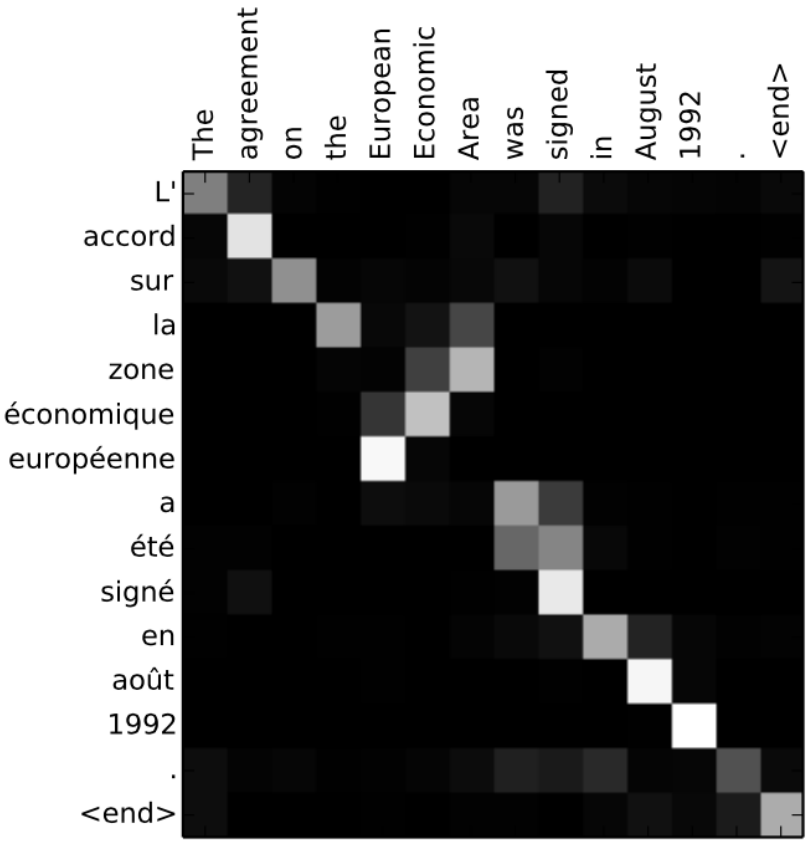

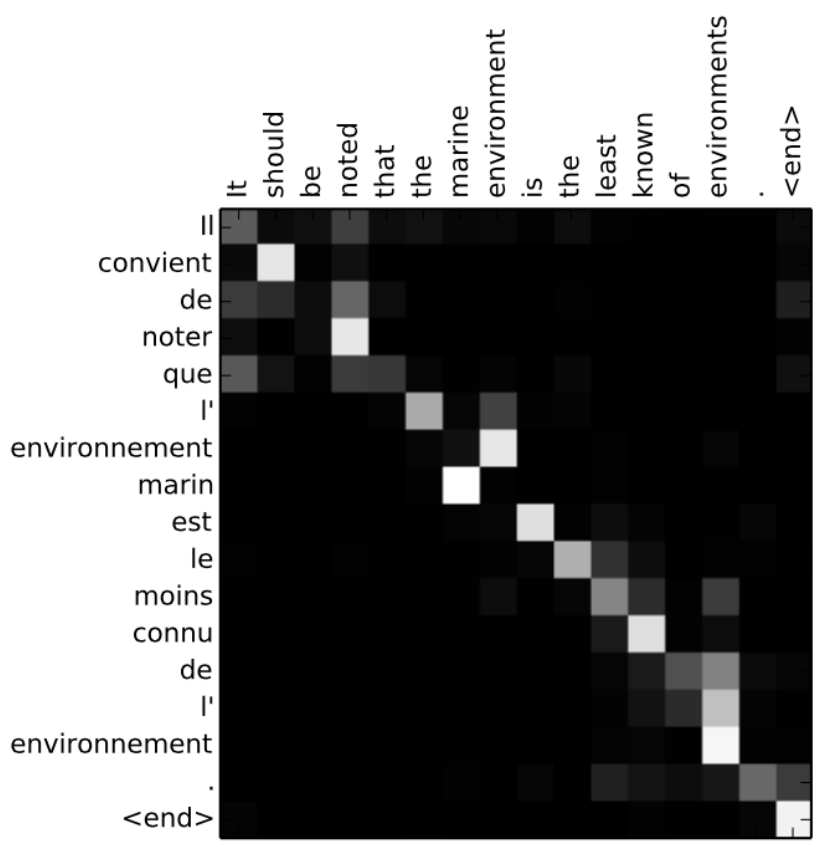

Sensible induced alignments

Attention-based translation

Bahdanau et al (2015)

Neural Machine Translation by Jointly Learning to Align and Translate

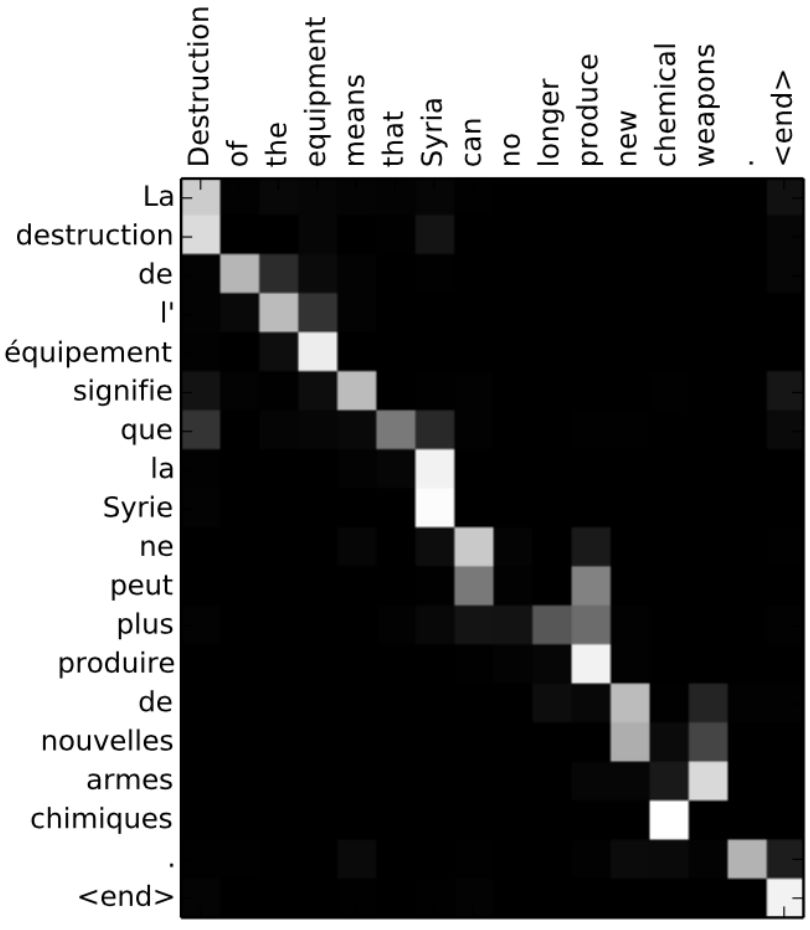

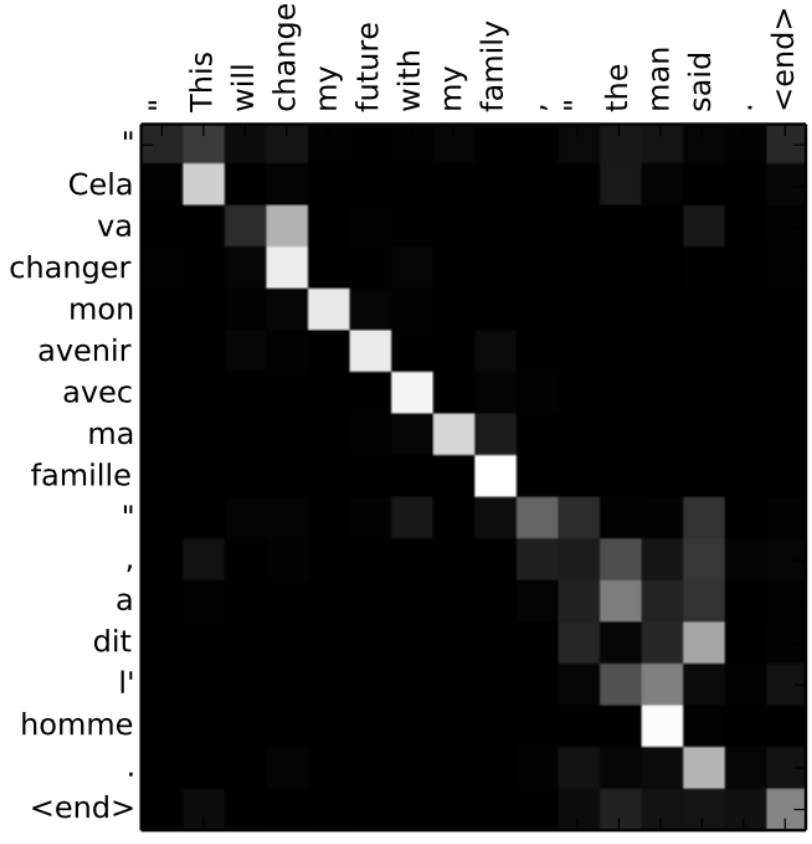

Sensible induced alignments

Natural language inference

Given a premise, e.g.,

> *The purchase of Houston-based LexCorp by BMI for $2Bn prompted widespread sell-offs by traders as they sought to minimize exposure. LexCorp had been an employee-owned concern since 2008.*

and a hypothesis, e.g.,

> *BMI acquired an American company.* (1)

predict whether the premise

- entails the hypothesis
- contradicts the hypothesis
- or remains neutral

Natural language inference

Given a premise, e.g.,

> *The purchase of Houston-based LexCorp by BMI for $2Bn prompted widespread sell-offs by traders as they sought to minimize exposure. LexCorp had been an employee-owned concern since 2008.*

and a hypothesis, e.g.,

> *BMI bought employee-owned LexCorp for $3.4Bn.* (2)

predict whether the premise

- entails the hypothesis
- contradicts the hypothesis
- or remains neutral

Natural language inference

Given a premise, e.g.,

> *The purchase of Houston-based LexCorp by BMI for $2Bn prompted widespread sell-offs by traders as they sought to minimize exposure. LexCorp had been an employee-owned concern since 2008.*

and a hypothesis, e.g.,

> *BMI is an employee-owned concern.* (3)

predict whether the premise

- entails the hypothesis
- contradicts the hypothesis
- or remains neutral

Natural language inference

Rocktäschel et al (2016)

Reasoning about Entailment with Neural Attention

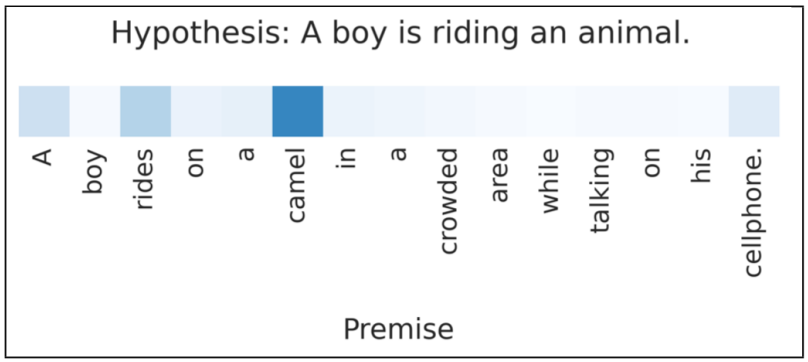

(a)

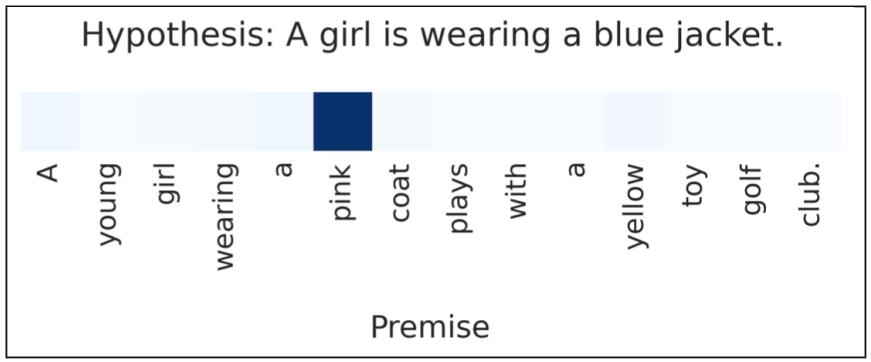

(b)

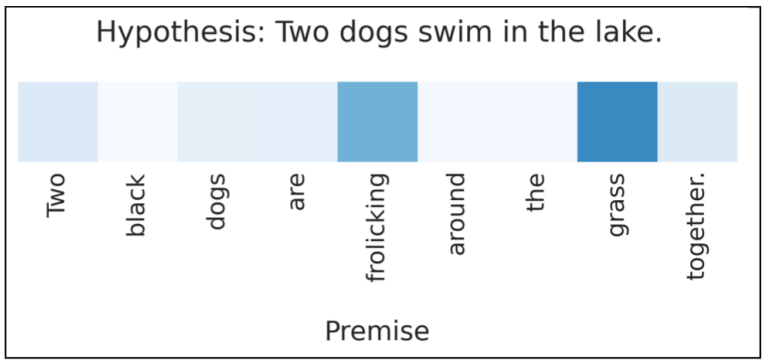

(c)

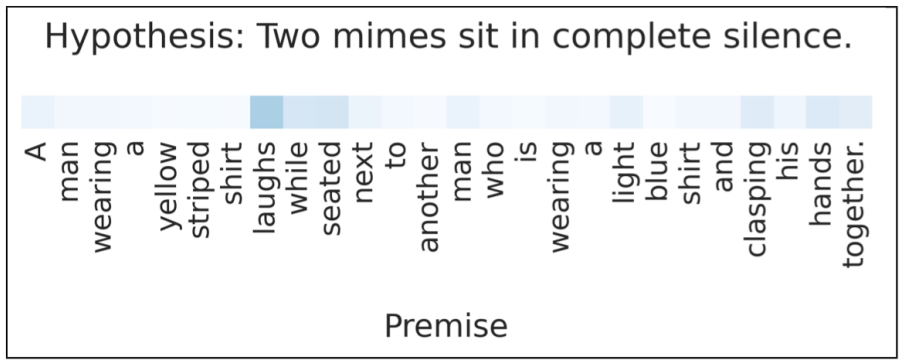

(d)

Attention conditioned on h_T

Natural language inference

Rocktäschel et al (2016)

Reasoning about Entailment with Neural Attention

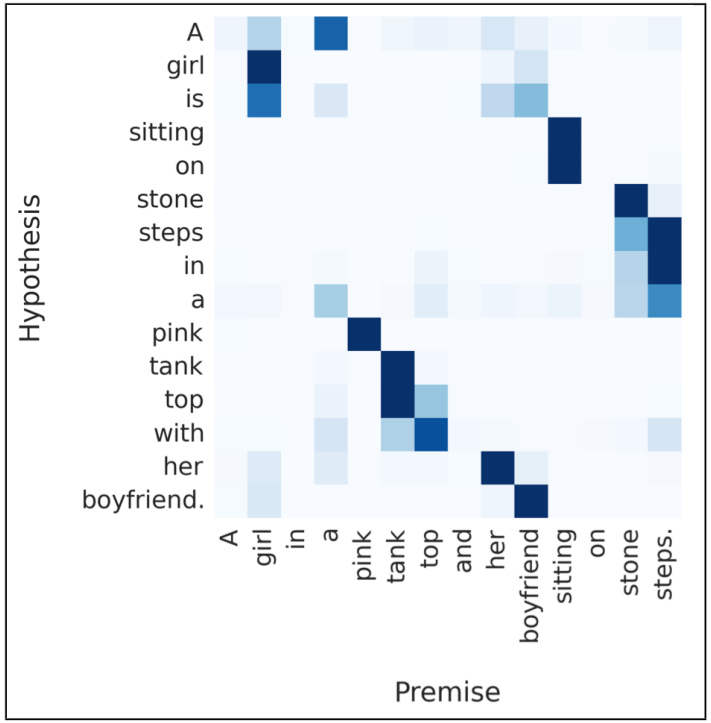

(a)

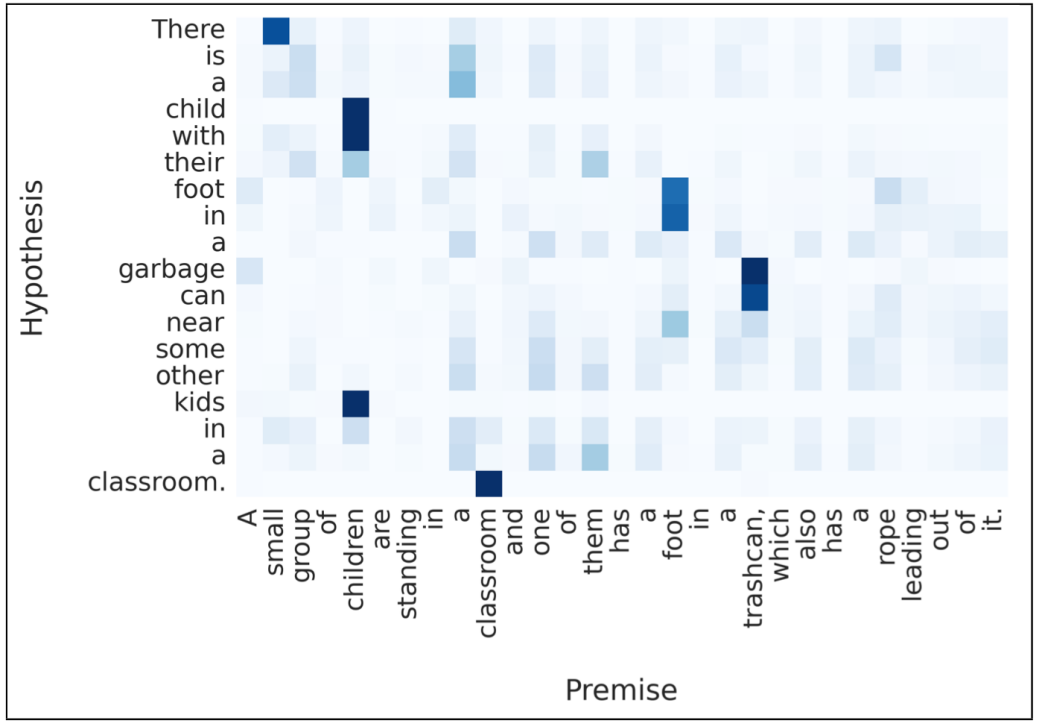

(b)

Attention conditioned on $h_1, \ldots, h_T$: Synonymy, importance

Natural language inference

Rocktäschel et al (2016)

Reasoning about Entailment with Neural Attention

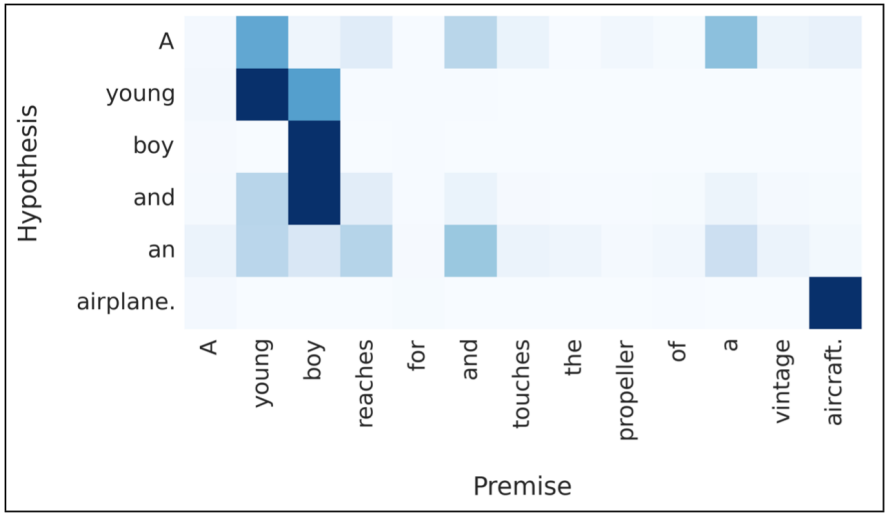

(c)

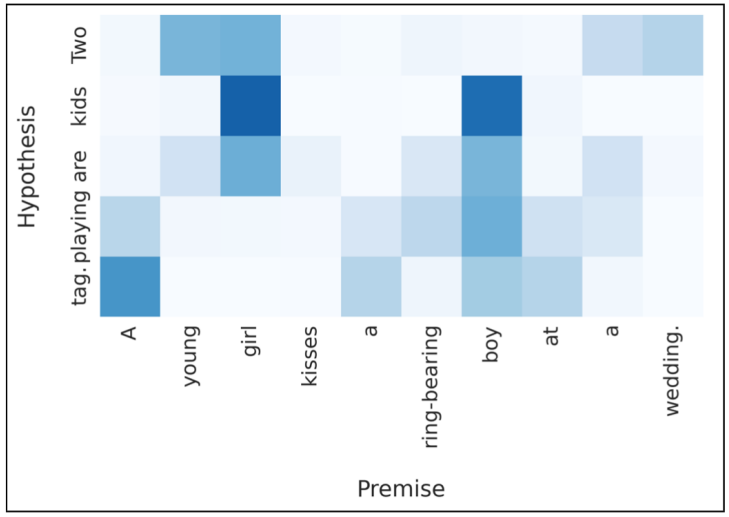

(d)

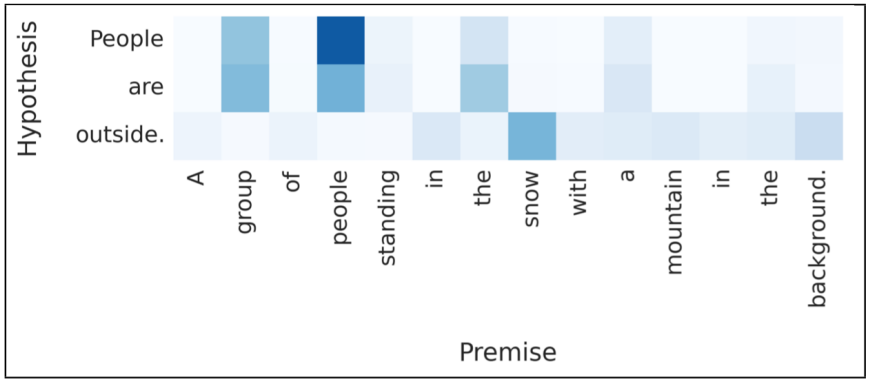

(e)

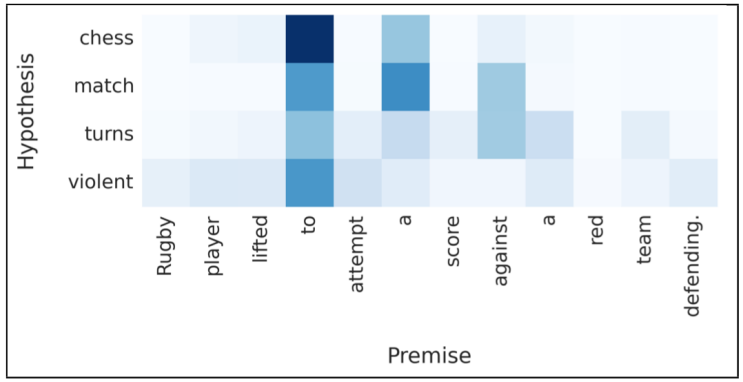

(f)

Attention conditioned on $h_1, \ldots, h_T$: Relatedness

Natural language inference

Rocktäschel et al (2016)

Reasoning about Entailment with Neural Attention

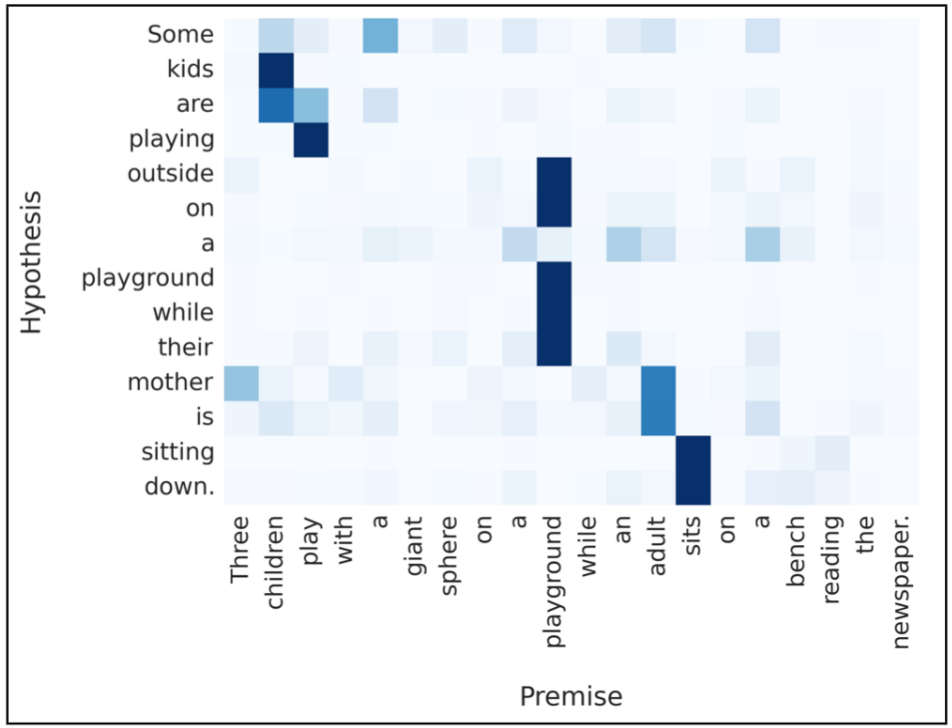

(g)

Attention conditioned on $h_1, \ldots, h_T$: Many:one

Self-attention

Cheng et al (2016)

Long Short-Term Memory-Networks for Machine Reading

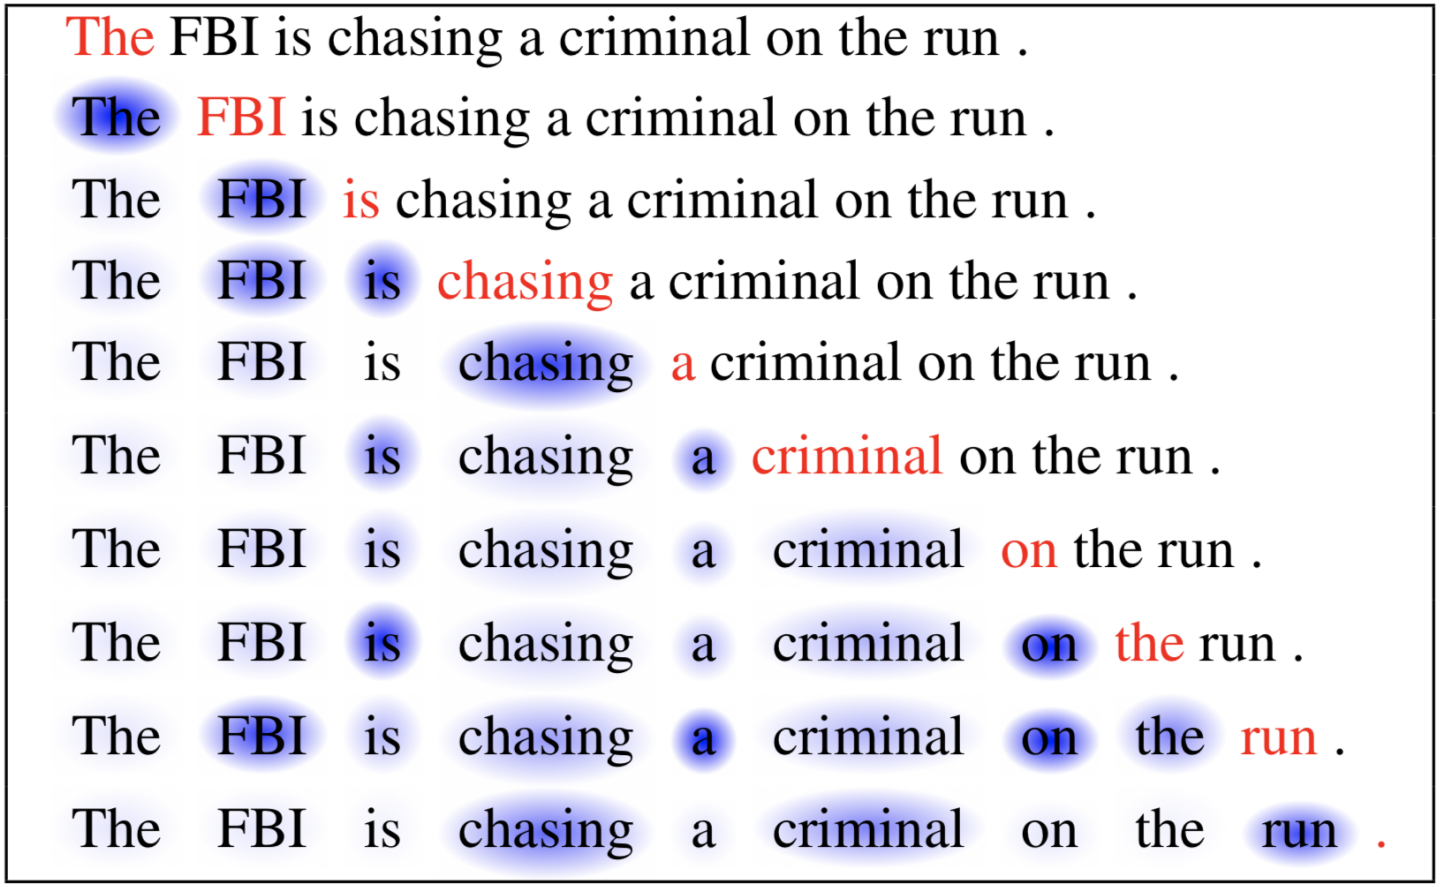

Attention over images

Xu et al (2015)

Show, Attend & Tell: Neural Image Caption Generation with Visual Attention

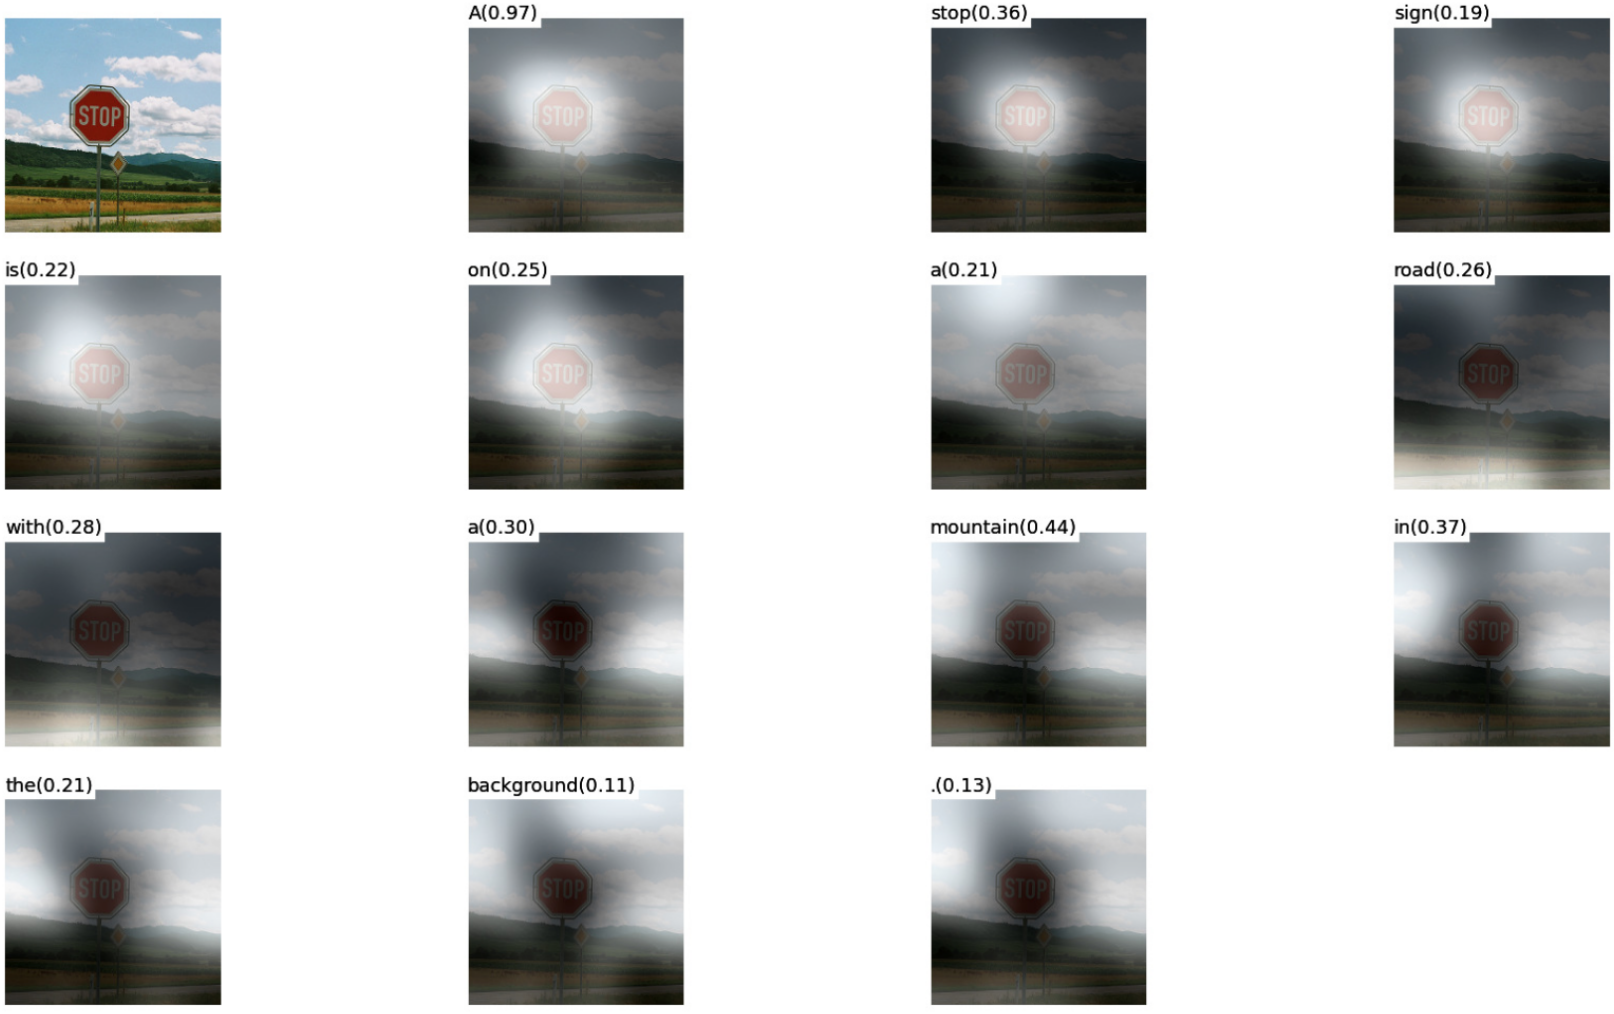

(b) A stop sign is on a road with a mountain in the background.

Attention over videos

Yao et al (2015)

Describing Videos by Exploiting Temporal Structure

***+Local+Global*: Someone is frying a fish in a pot**

Attention variants

$c = \textsc{Attention}(\text{query } q, \text{ keys } k_1 \dots k_n)$

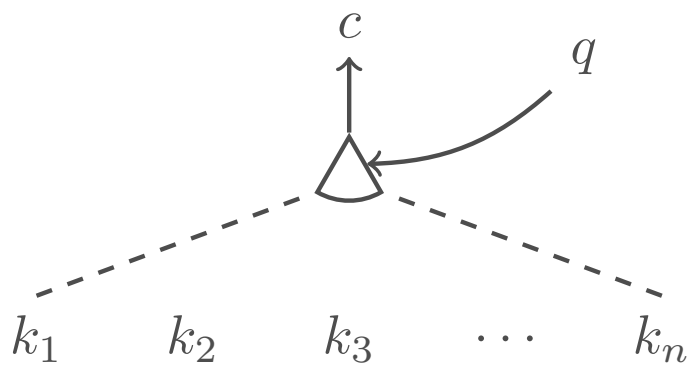

$$\alpha_i = \text{softmax}(\text{score}(q, k_i))$$

$$c = \sum_i \alpha_i, k_i$$

Attention variants

$c = \textsc{Attention}(\text{query } q, \text{ keys } k_1 \dots k_n, \text{ values } v_1 \dots v_n)$

e.g., memory networks (Weston et al, 2015; Sukhbataar et al, 2015)

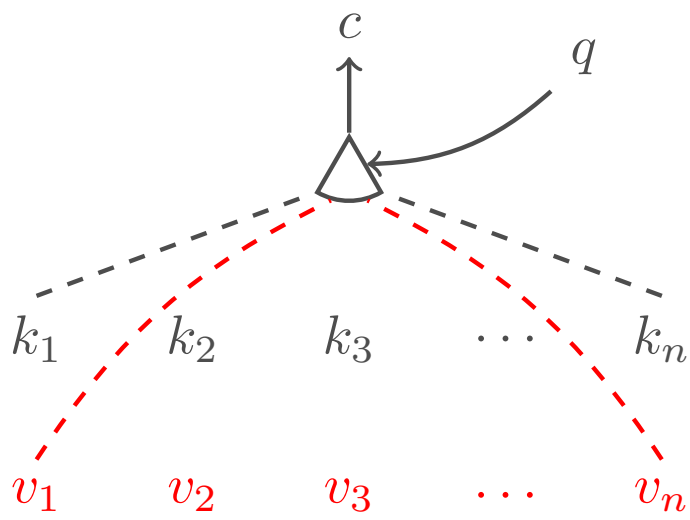

$$\alpha_i = \text{softmax}(\text{score}(q, k_i))$$

$$c = \sum_i \alpha_i, v_i$$

Attention variants

Weston et al (2015)

Memory Networks

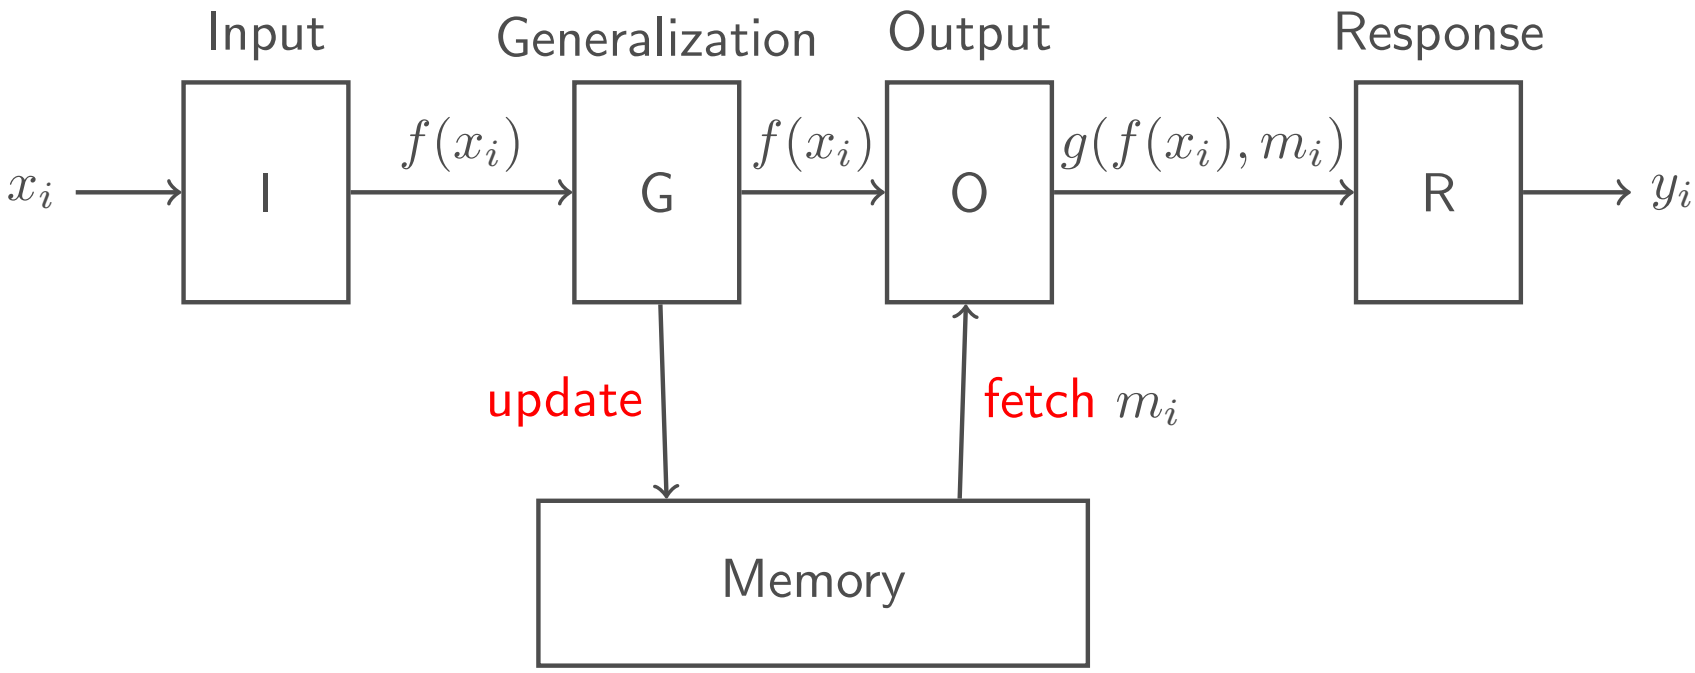

MemN2N (Sukhbataar et al, 2015)

+ Soft attention over memories
+ Multiple memory lookups (hops)
+ End-to-end training

Attention scoring functions

- Additive (Bahdanau et al, 2015)

$$\text{score}(q, k) = u^\top \tanh(W[q; k])$$

- Multiplicative (Luong et al, 2015)

$$\text{score}(q, k) = q^\top W k$$

- Scaled dot-product (Vaswani et al, 2017)

$$\text{score}(q, k) = \frac{q^\top k}{\sqrt{d_k}}$$

Attention variants

- Stochastic hard attention (Xu et al, 2015)
- Local attention (Luong et al, 2015)
- Monotonic attention (Yu et al, 2016; Raffel et al, 2017)
- Self attention (Cheng et al, 2016; Vaswani et al, 2017)
- Convolutional attention (Allamanis et al, 2016)
- Structured attention (Kim et al, 2017)
- Multi-headed attention (Vaswani et al, 2017)

Transformer

Vaswani et al (2017)

Attention is All You Need

RNN encoder

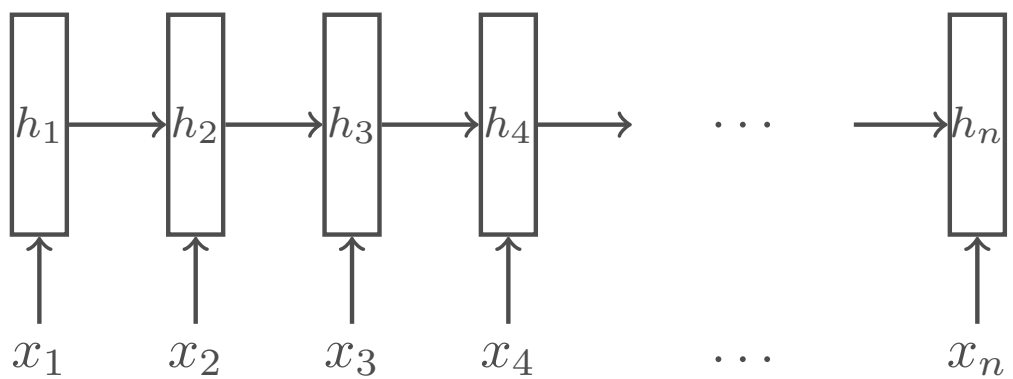

Transformer

Vaswani et al (2017)

Attention is All You Need

RNN encoder with attention

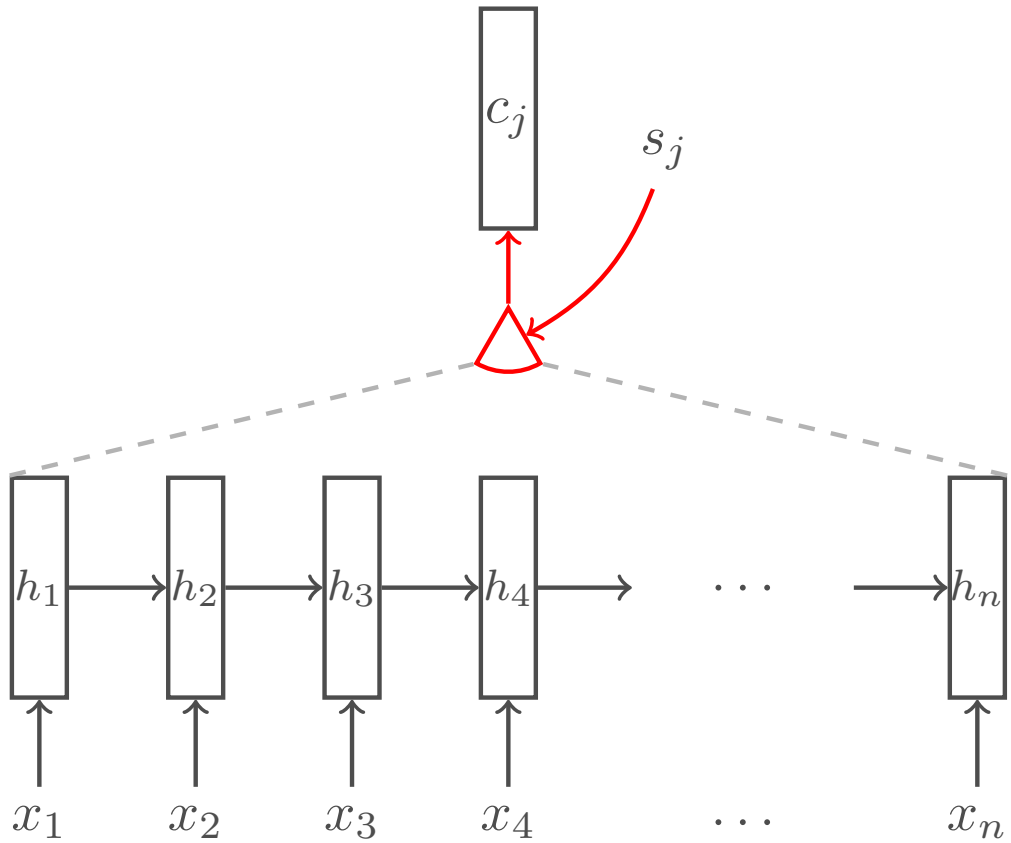

Transformer

Vaswani et al (2017)

Attention is All You Need

Deep encoder with self-attention

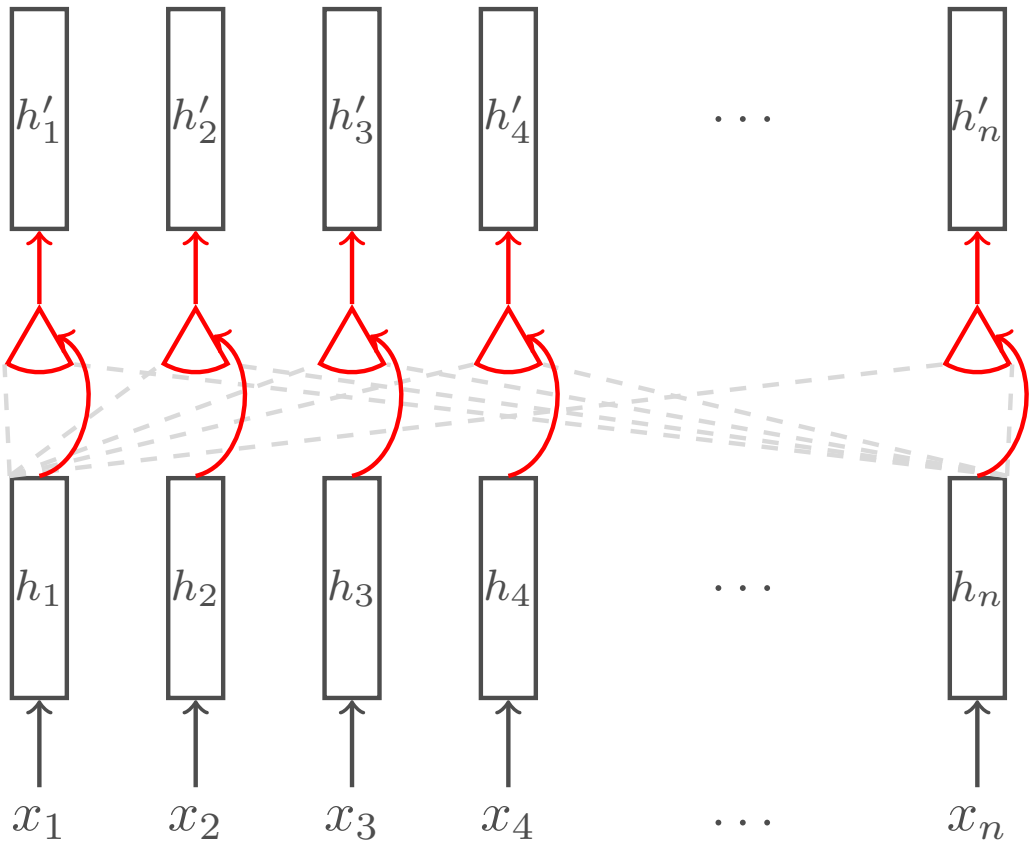

Transformer

Vaswani et al (2017)

Attention is All You Need

Deep encoder with multi-headed self-attention

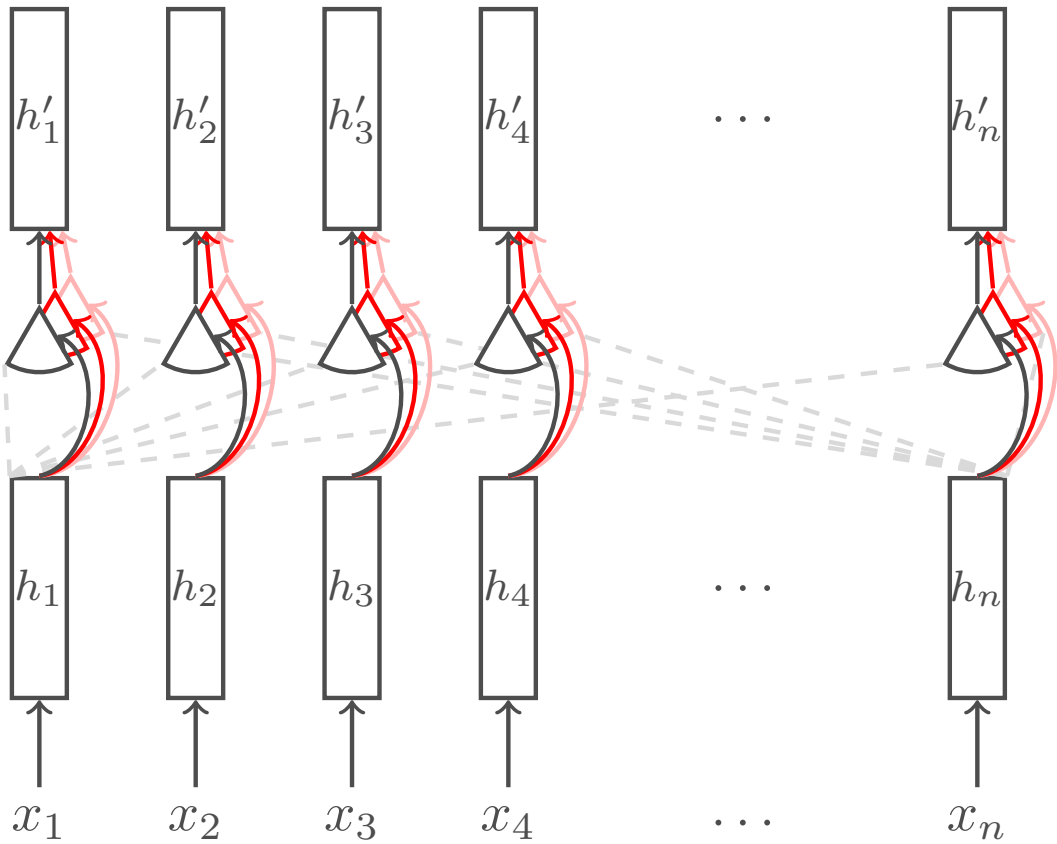

Transformer

Vaswani et al (2017)

Attention is All You Need

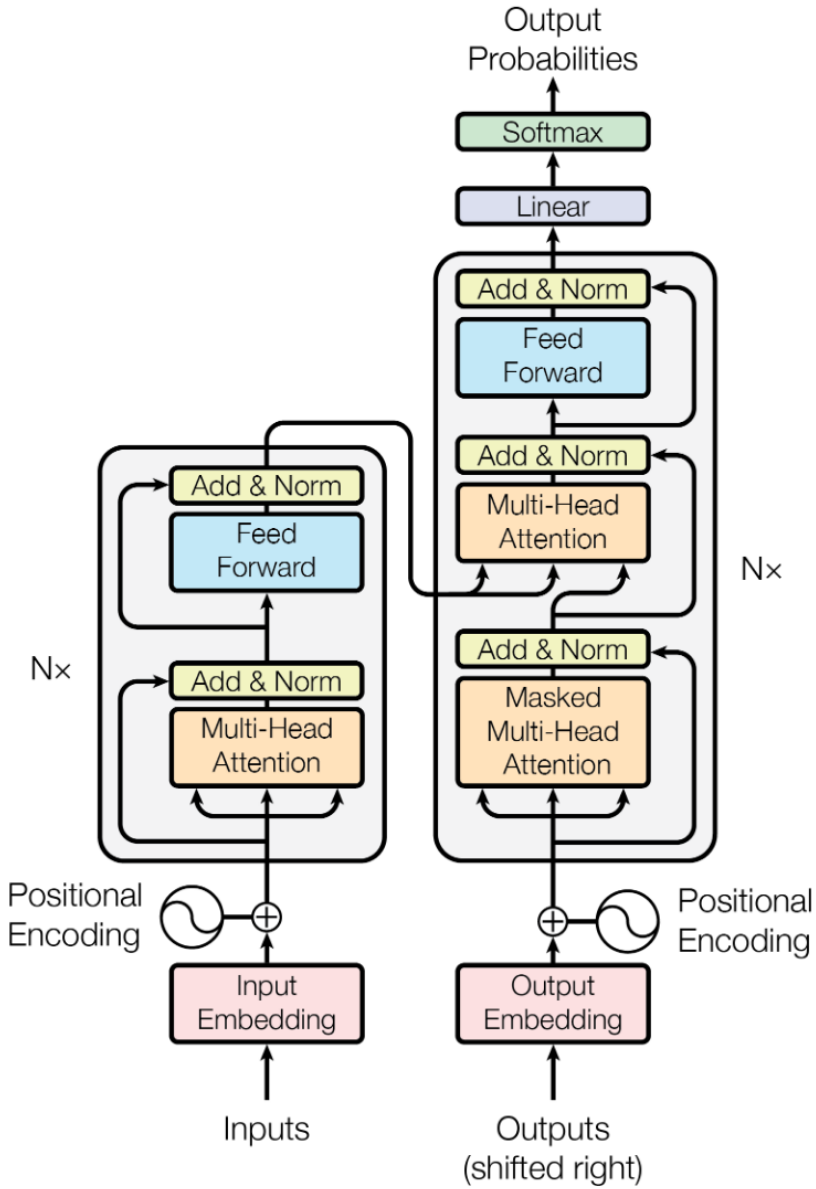

Transformer

Vaswani et al (2017)

Attention is All You Need

- Self-attention at every layer instead of recurrence
  - \+ Inference can be parallelized
- No sensitivity to input position
  - – Positional embeddings required
  - \+ Can apply to sets
- Deep architecture (6 layers) with multi-headed attention
  - \+ Higher layers appear to learn linguistic structure
- Scaled dot-product attention with masking
  - \+ Avoids bias in simple dot-product attention
  - \+ Fewer parameters needed for rich model
- Improved runtime and performance on translation, parsing, etc

Transformer

Vaswani et al (2017)

Attention is All You Need

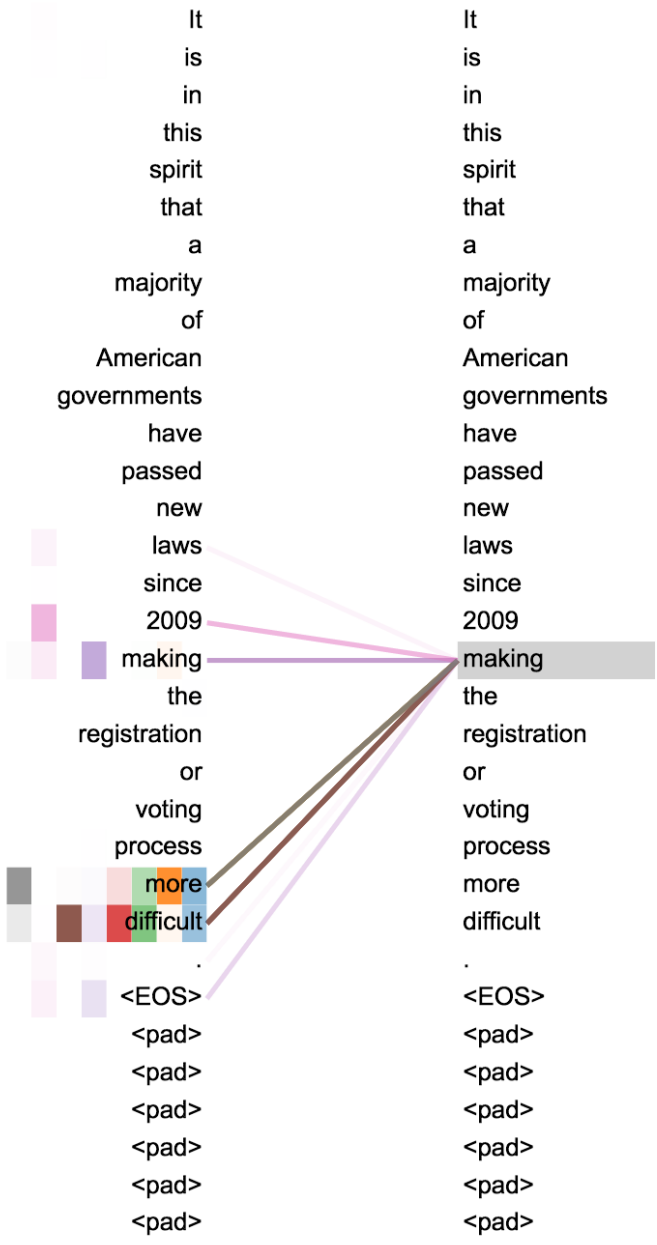

Transformer

Vaswani et al (2017)

Attention is All You Need

Transformer

Vaswani et al (2017)

Attention is All You Need

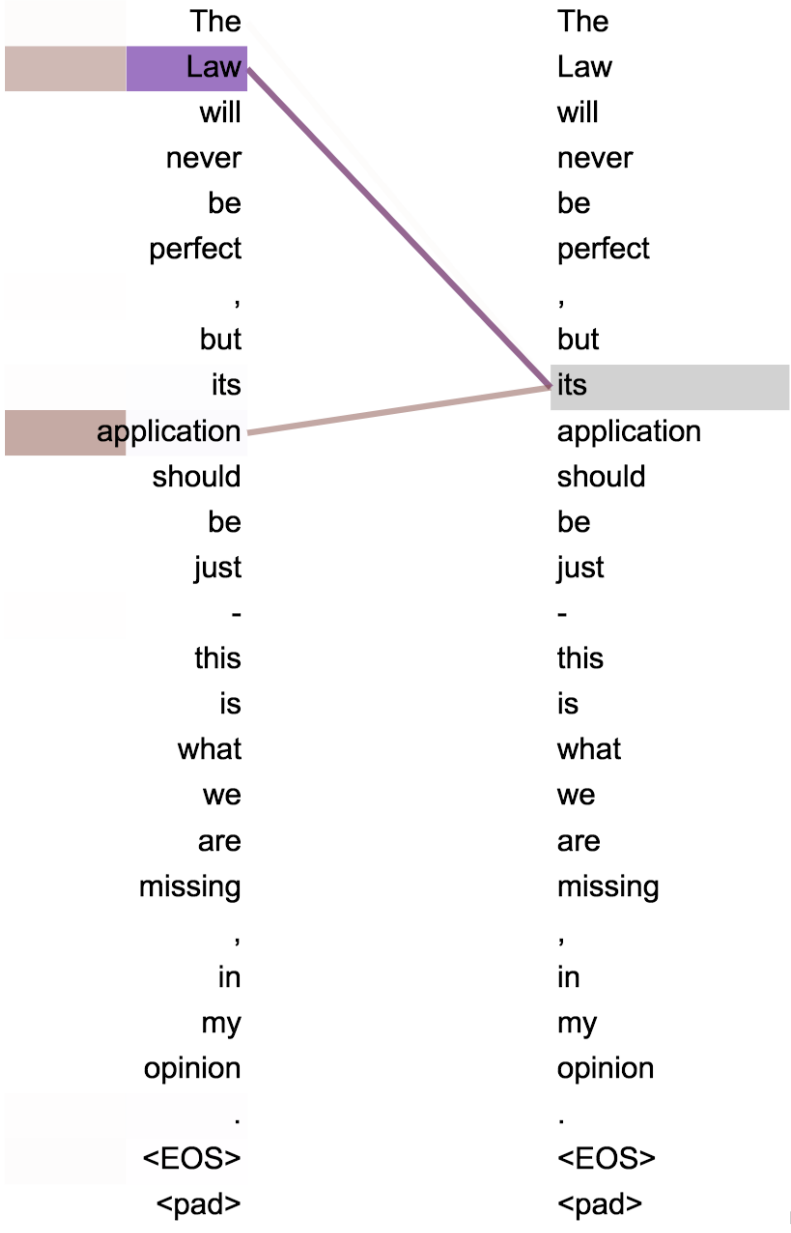

Transformer

Vaswani et al (2017)

Attention is All You Need

Transformer

Vaswani et al (2017)

Attention is All You Need

Large vocabularies

Sequence-to-sequence models can typically scale to 30K-50K words

But real-world applications need at least 500K-1M words

Large vocabularies

Alternative 1: Hierarchical softmax

- Predict path in binary tree representation of output layer
- Reduces to $\log_2(V)$ binary decisions

$$p(w_t = \text{“dog”} | \cdots) = (1 - \sigma(U_0 h_t)) \times \sigma(U_1 h_t) \times \sigma(U_4 h_t)$$

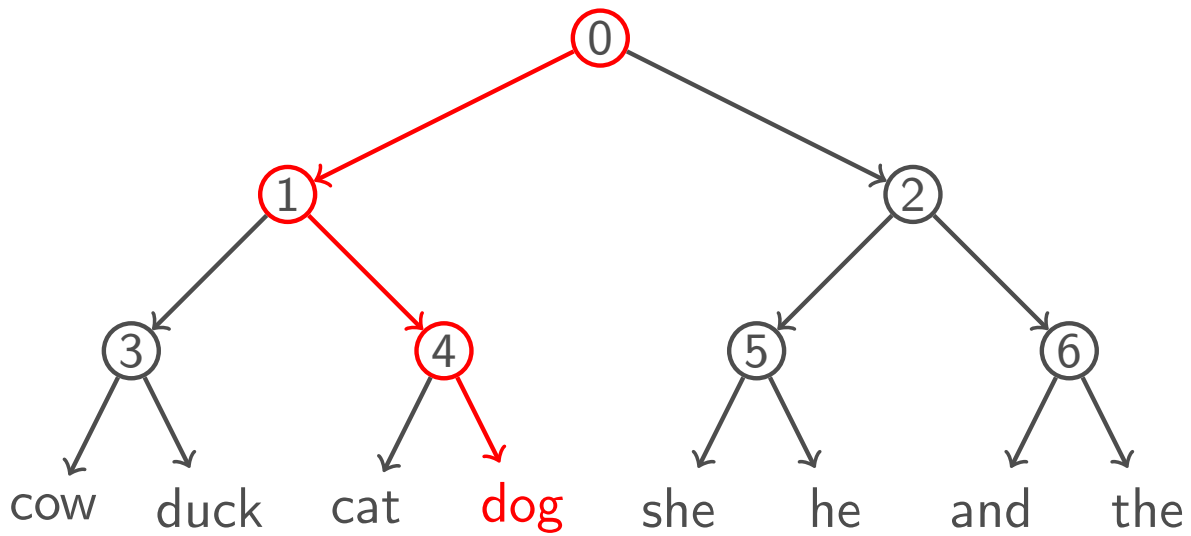

Large vocabularies

Jean et al (2015)

On Using Very Large Target Vocabulary for Neural Machine Translation

Alternative 2: Importance sampling

- Expensive to compute the softmax normalization term over V

$$p(y_i = w_j | y_{<i}, x) = \frac{\exp\left(W_j^\top f(s_i, y_{i-1}, c_i)\right)}{\sum_{k=1}^{|V|} \exp\left(W_k^\top f(s_i, y_{i-1}, c_i)\right)}$$

- Use a small subset of the target vocabulary for each update

- Approximate expectation over gradient of loss with fewer samples

- Partition the training corpus and maintain local vocabularies in each partition to use GPUs efficiently

Large vocabularies

Sennrich et al (2016)

Neural machine translation of rare words with subword units

Alternative 3: Subword units

- Reduce vocabulary by replacing infrequent words with sub-words

Jet makers feud over seat width with big orders at stake

⇓

_J et _makers _fe ud _over _seat _width _with _big _orders _at _stake

- Code for byte-pair encoding:
  `https://github.com/rsennrich/subword-nmt`

Copying mechanism

Gu et al (2016)

Incorporating Copying Mechanism in Sequence-to-Sequence Learning

In monolingual tasks, copy rare words directly from the input

- Generation via standard attention-based decoder

$$\psi_g(y_i = w_j) = W_j^\top f(s_i, y_{i-1}, c_i) \qquad w_j \in V$$

- Copying via a non-linear projection of input hidden states

$$\psi_c(y_i = x_j) = \tanh(h_j^\top U) f(s_i, y_{i-1}, c_i) \qquad x_j \in X$$

- Both modes compete via the softmax

$$p(y_i = w_j | y_{<i}, x) = \frac{1}{Z} \left(\exp\left(\psi_g(w_j)\right) + \sum_{k:x_k = w_j} \exp\left(\psi_c(x_k)\right) \right)$$

Copying mechanism

Gu et al (2016)

Incorporating Copying Mechanism in Sequence-to-Sequence Learning

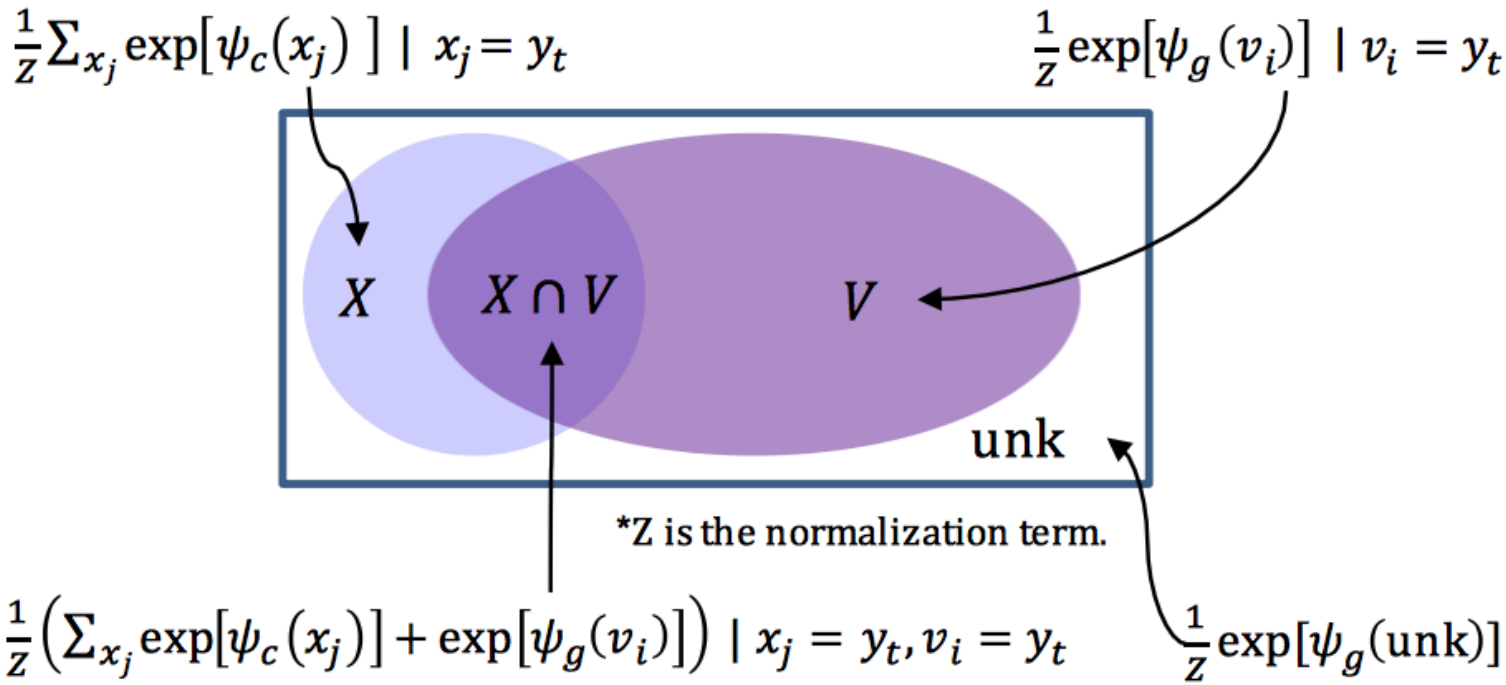

Decoding probability $p(y_t|\cdots)$

Copying mechanism

Gu et al (2016)

Incorporating Copying Mechanism in Sequence-to-Sequence Learning

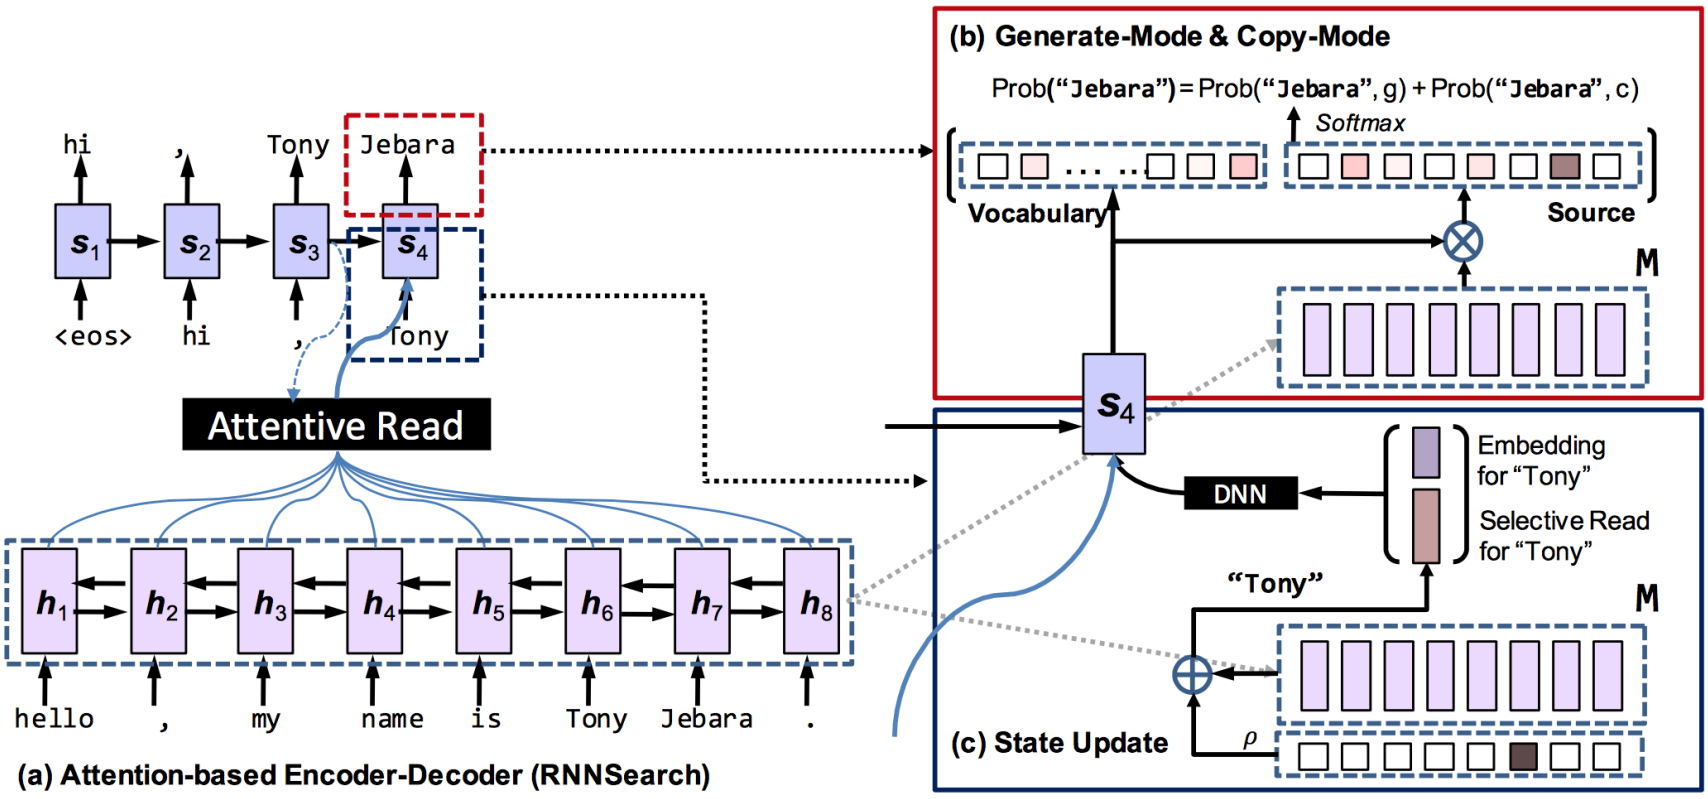

Copying mechanism

See et al (2017)

Get to the Point: Summarization with Pointer Generator Networks

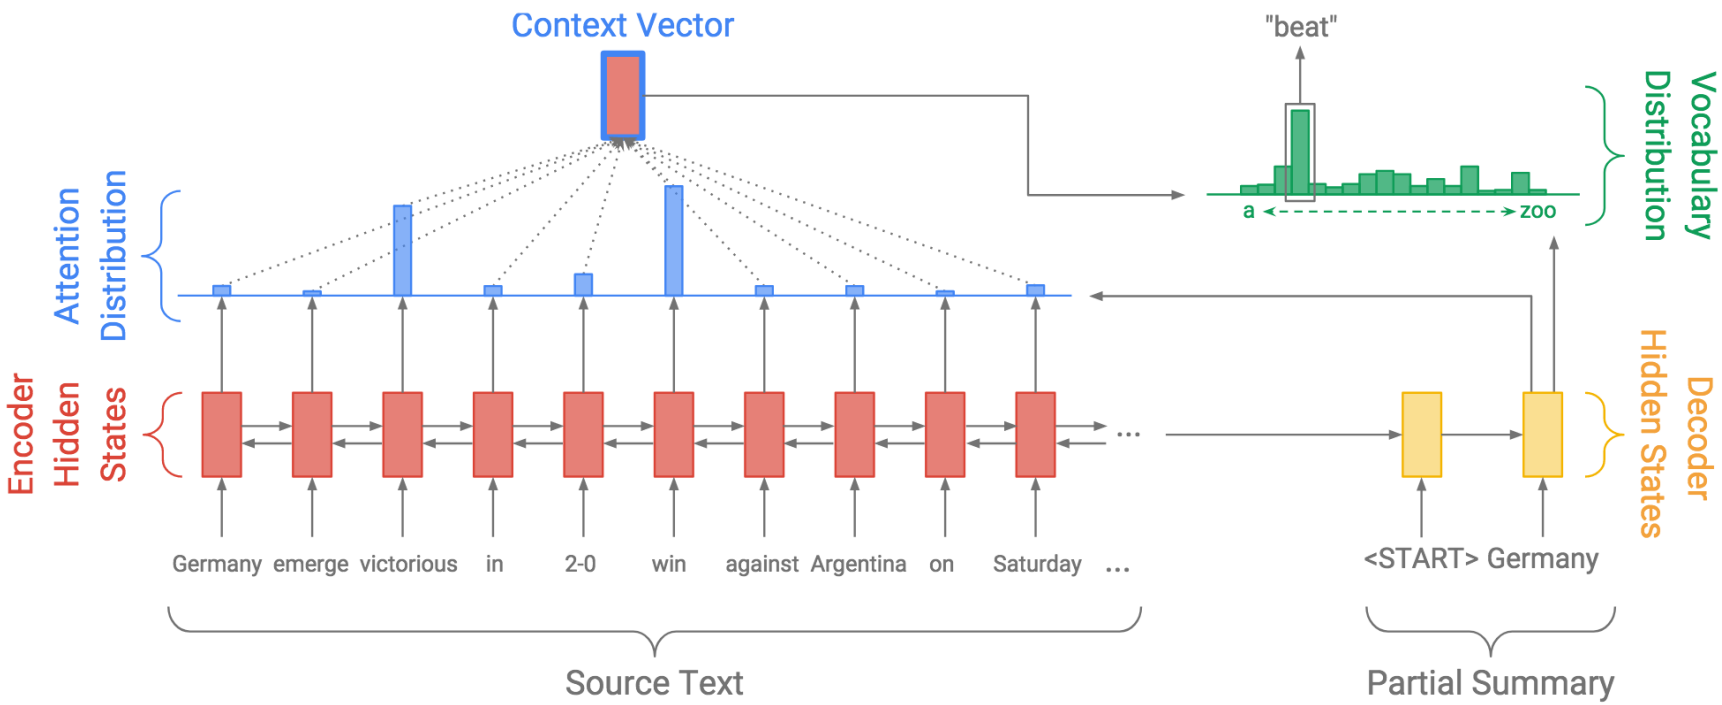

Attention for common words

Copying mechanism

See et al (2017)

Get to the Point: Summarization with Pointer Generator Networks

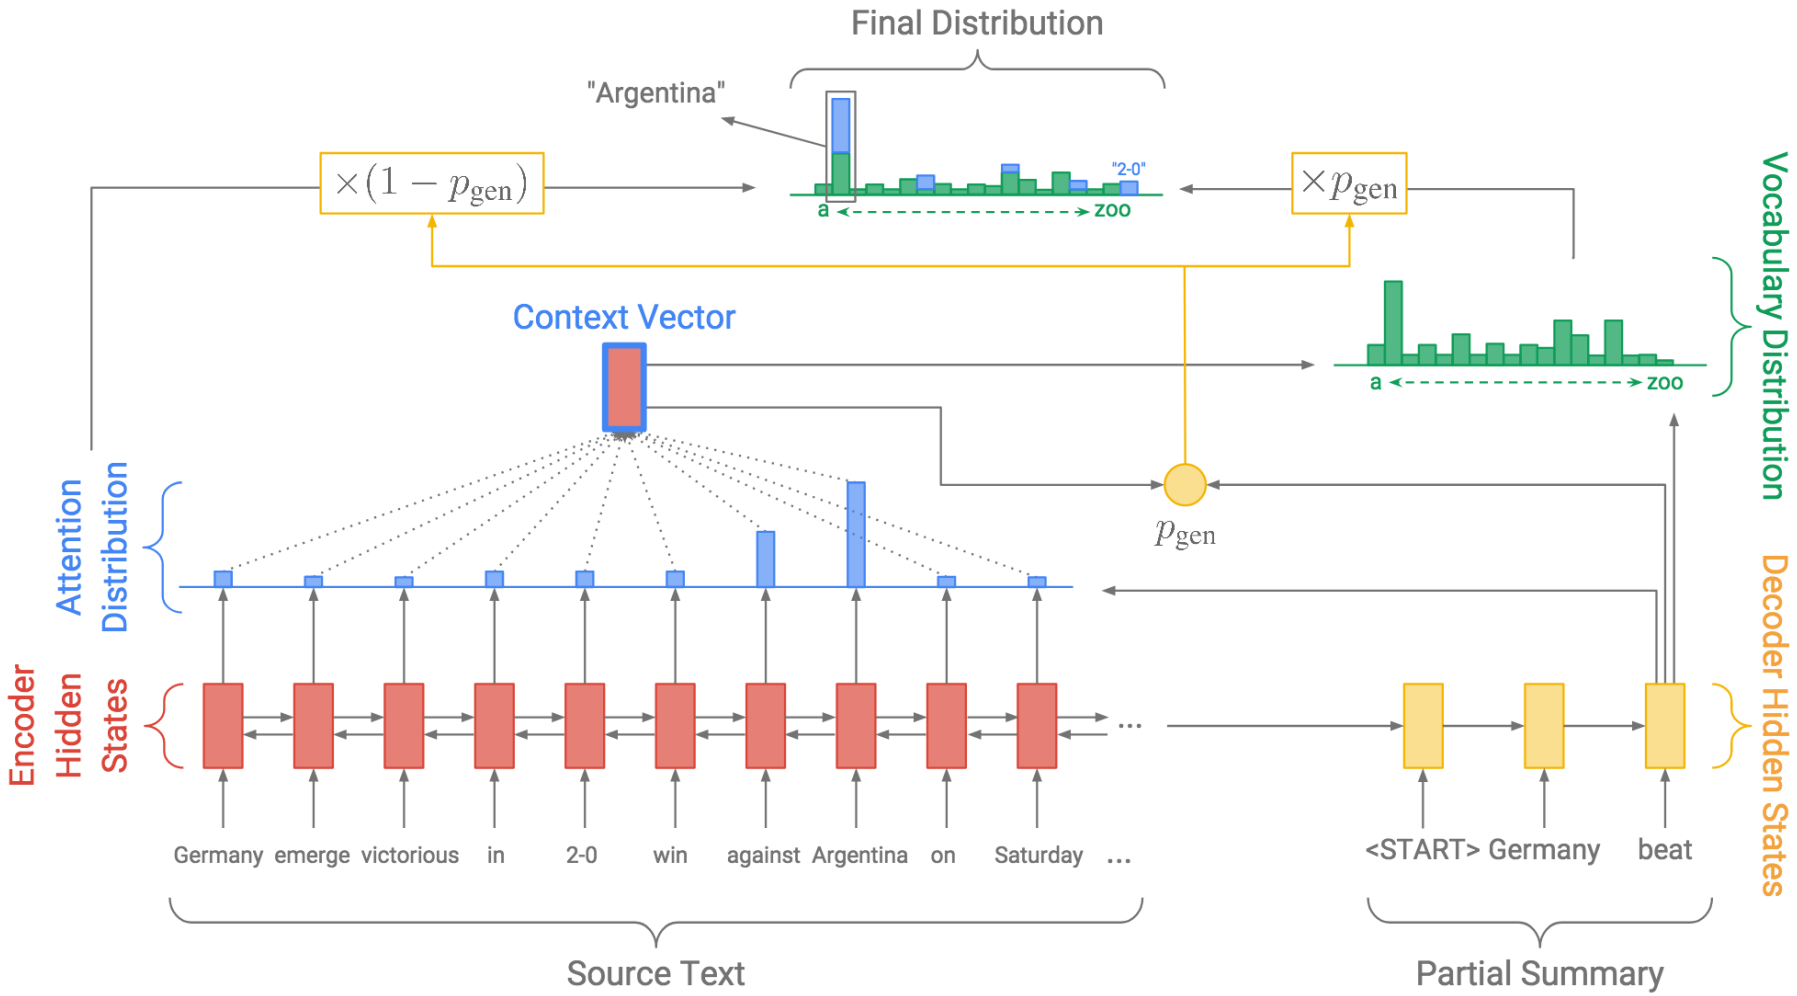

Copying from input for rarer words

Autoencoders

Given input x, learn an encoding z that can be decoded to reconstruct x

For sequence input $x_1, \ldots, x_n$, can use standard MT models

- Is attention viable?

\+ Useful for pre-training text classifiers (Dai et al, 2015)

Denoising autoencoders

Hill et al (2016)

Learning Distributed Representations of Sentences from Unlabeled Data

Given noisy input $\tilde{x}$, learn an encoding z that can be decoded to reconstruct x

Noise functions for sequences

- Drop words with probability p_{drop}
- Swap words with probability p_{swap}

\+ Useful as features for a linear classifier

\+ Can learn sentence representations without ordered docs

Dataset		Sentence 1	Sentence 2	/5
	News	*Mexico wishes to guarantee citizens' safety.*	*Mexico wishes to avoid more violence.*	4
	Forum	*The problem is simpler than that.*	*The problem is simple.*	3.8
STS	WordNet	*A social set or clique of friends.*	*An unofficial association of people or groups.*	3.6
2014	Twitter	*Taking Aim #Stopgunviolence #Congress #NRA*	*Obama, Gun Policy and the N.R.A.*	1.6
	Images	*A woman riding a brown horse.*	*A young girl riding a brown horse.*	4.4
	Headlines	*Iranians Vote in Presidential Election.*	*Keita Wins Mali Presidential Election.*	0.4
SICK (test+train)		*A lone biker is jumping in the air.*	*A man is jumping into a full pool.*	1.7

Variational autoencoders (VAEs)

Kingma & Welling (2015)

Auto-encoding Variational Bayes

Autoencoders often don't generalize well to new data, noisy representations

Approximate the posterior $p(z|x)$ with variational inference

- Encoder: induce $q(z|x)$ with parameters θ
- Decoder: sample z and reconstruct x with parameters ϕ
- Loss:

$$\ell_i = -\mathbb{E}_{z \sim q_\theta(z|x_i)} \log p_\phi(x_i|z) + \mathsf{KL}(q_\theta(z|x_i)||p(z))$$

Estimate gradients using *reparameterization trick* for Gaussians

$$\mathcal{N}(\mu, \sigma^2) = \mathcal{N}(0, 1)\sigma + \mu$$

\+ Can interpolate smoothly in representation space, e.g., `https://giphy.com/gifs/26ufgj5LH3YKO1Zlu`