

Starting from ImageNet Model: Fine-Tuning and Adversarial Attack

Liangliang Cao

<https://columbia6894.github.io/>

Administrative

By this week you should:

1. form final project team (3 ppl per team)
2. decide the topic of your final projects

Then you can

- finish 1st problem in your homework 3
- receive google cloud credits (by contacting TA)

Administrative (Continued)

Have problems with your final project? Try the following

1. Choose a **recent** paper that you are interested in
2. Implement it or test it by yourself
3. Then you can
 - a) If it fails in some scenario, try to improve it
 - b) If it works great, apply it to a new application

For more details, refer to

http://llcao.net/cu-deeplearning17/lecture/lecture5_llc.pdf

Outline

Last vision lecture discussed learning from ImageNet/Celeb 1M

You can impress others with a model pretrained from ImageNet:

- Fine-tuning
- Adversarial attack

Fine Tuning Models from ImageNet

Models trained from ImageNet have learned effective feature presentation.

We can treat deep CNNs as feature extractors, and fine tune a new model over it.

Several ways to do the fine tuning:

1. Train Linear SVM over deep features (least training examples)
2. Fine tuning only the cross-entropy (more training examples)
3. Fine tuning both classifier and deep CNNs (about thousands of examples)

1. Fine-tuning using linear SVM

Recommend Liblinear for linear SVM:

<https://www.csie.ntu.edu.tw/~cjlin/liblinear/>

Steps:

1. Take the last layer of ImageNet model as feature extractor
2. Extract the feature for all the training examples
3. Optimize liblinear model and fine the global minimization

2. Fine-tuning the top layers

Steps:

1. Freeze the convolutional layers of pre-trained model
 `model = applications.VGG16(weights='imagenet', include_top=False)`
 for layer in model.layers[:25]:
 `layer.trainable = False`
2. Add a top model for the new classification task
3. Train the top model using SGD optimization

One easy-to-follow reference: <https://towardsdatascience.com/a-comprehensive-guide-on-how-to-fine-tune-deep-neural-networks-using-keras-on-google-colab-free-daaaa0aced8f>

Differences

	SVM	New Top layer
Extra toolkit	Liblinear	Keras/Tensorflow
Optimization	Global minimum	SGD may fall to local minimum
Number of training	A few images to dozens	Dozens to hundreds

Tricks of improve fine-tuning:

1. For SGD, use slow training rate first, do not use RMSProp
2. Data augmentation may be very helpful
3. Regularization may help fine-tuning (though we do not do regularization in general deep learning.)

3. One Step Further

When there are more training samples, we shall consider fine-tune the convolutional filters as well:

- Use smaller learning rates for CNN filters but bigger rates for the top layer.
- Monitor the training loss whether training accuracy = 100%
- Reference:
http://caffe.berkeleyvision.org/gathered/examples/finetune_flickr_style.html

Future research questions you may consider for final projects:

- Actively choose samples to label? (with a budget)
- Efficiently learn new categories, e.g., for face recognition

Adversarial Attack

How many of you think deep CNNs are very reliable?

Guess what an ImageNet model will predict:



Adversarial Attack


How many of you think deep CNNs are very reliable?

Guess what an ImageNet model will predict:




Adversarial Attack


African elephant




koala



Difference



10x Difference




schooner




iPod



Difference



10x Difference



Examples are courtesy to Stanford cs231n lecture slides.

How to Compute the Adversarial Example

$$J(\tilde{\mathbf{x}}, \boldsymbol{\theta}) \approx J(\mathbf{x}, \boldsymbol{\theta}) + (\tilde{\mathbf{x}} - \mathbf{x})^\top \nabla_{\mathbf{x}} J(\mathbf{x})$$

Maximize $(\tilde{\mathbf{x}} - \mathbf{x})^\top \nabla_{\mathbf{x}} J(\mathbf{x})$

Subject to $\|\tilde{\mathbf{x}} - \mathbf{x}\|_\infty \leq \epsilon$

The Fast Gradient
Sign Method

So the adversarial example can be generated by

$$\tilde{\mathbf{x}} = \mathbf{x} + \epsilon \text{sign}(\nabla_{\mathbf{x}} J(\mathbf{x}))$$

How to Implement?

There are a number of toolboxes such as cleverhans, Foolbox, etc

```
from cleverhans.attacks import FastGradientMethod
fgsm = FastGradientMethod(model, sess=sess)
fgsm_params = {'eps': 0.3, 'clip_min': 0., 'clip_max': 1.}
adv_x = fgsm.generate_np(origin_x, **fgsm_params)
```

But fundamentally it is just to compute the gradient subject to input x .
You should read the code of cleverhans or Foolbox by yourself.

From White-box Attack to Black-box Attack

Fast Gradient Sign Method (fgsm) requires to know the model parameters to compute the adversarial attack. It is called a white-box attack coz we know the details of the model.

In practice attacker does not know the model. They can

- evaluate the model multiple times to approximate the gradient
- attack some venerable tasks such as object detection or QA

For final projects you may refer to:

- [NIPS 2017 adversarial attack competition](#)
- [Percy Liang](#) and [Dawn Song](#)'s work on adversarial attack systems
- [The Elephant in the Room](#) attack for object detection